

Lecture 8: Linear Classifiers and More Model Validation

INFO 1998: Introduction to Machine Learning



CDS Education

We explore, learn, and educate big minds.

Agenda

1. **Perceptron + SVM**
2. **More Cross-Validation techniques**




Linear Classifiers



Linear Classifiers

A linear classifier is a hyper plane that is used to classify our data points

A hyperplane is our **decision boundary** and our goal is to find the hyper plane that best classifies our data



Perceptron Learning Algorithm

Goal: find a normal vector w that perfectly classifies all the points in our data set

Algorithm:

Initialize classifier as some random hyperplane

While there exists a misclassified point x :


 Tilt classifier slightly so that it classifies x correctly
 (or, is a little closer to classifying x correctly)

End While

“Use your mistakes as your stepping stones”



#0



Also, Frank Rosenblatt was first to implement perceptron

Gave him the title of 'Father of Deep Learning'

He went to Cornell!!!



Limitations of Perceptron



Is a great model to understand the intuition behind the training of a linear classifier: iteratively improve classifier by using misclassified points 😊

The training algorithm will never terminate if your training dataset is not linearly separable 😞



Not Linearly Separable


This data set is not linearly separable because of an outlier




SVM




Classify (+) and (-)



Which Hyperplane?




Optimal Hyperplane




Maximal Margin Classifier

- We want to find a **separating hyperplane**
- Once we find candidates for the hyperplane, we try to maximize the **margin**, the normal distance from borderline points
 - Only **Support Vectors** matter




What if...




Which Decision Boundary is better?

Boundary 1



Boundary 2




Margins

Use cost function to penalize misclassified points

Choice of cost function makes margin “hard” vs. “soft”

Non-separable training sets

Use linear separation, but admit training errors.




Penalty of error: distance to hyperplane multiplied by *error cost* C .



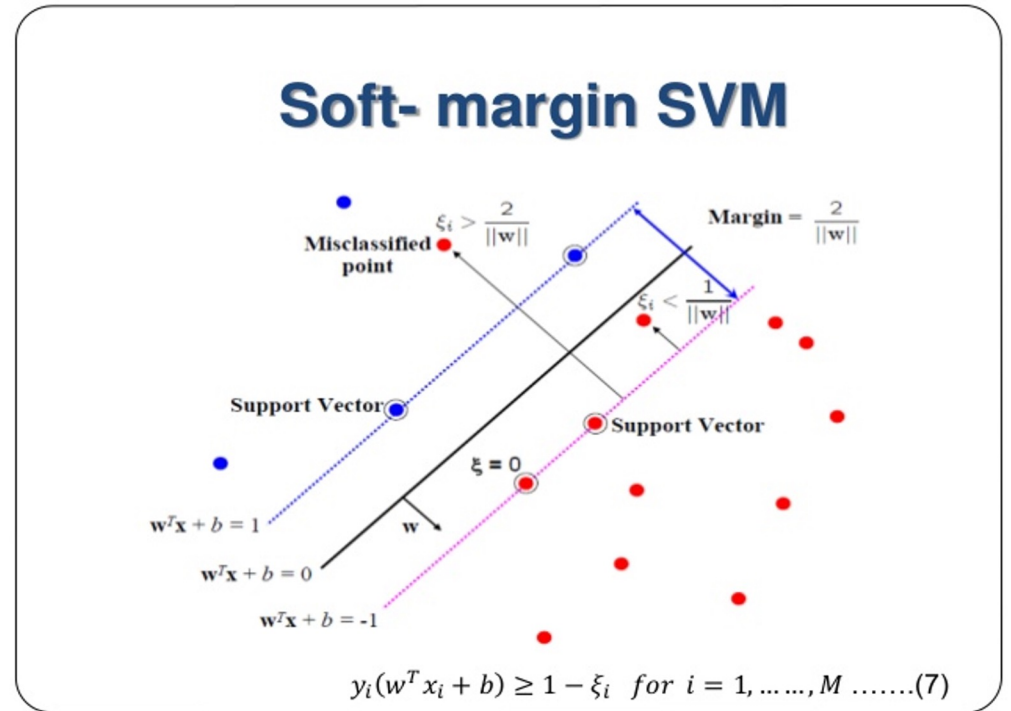
Hard Margins

- High penalty value
- The hyperplane can be dictated by a single outlier

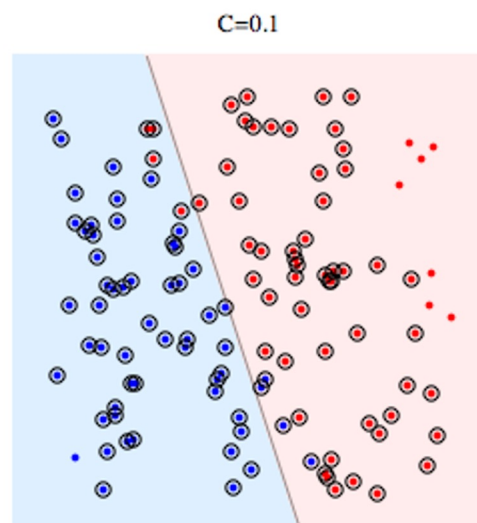




Soft Margins

- Used in non-linearly separable datasets
- Allow for misclassification
- Can account for “dirty” boundaries




Misclassification Penalty C





Kernels

- You cannot linearly divide the 2 classes on the xy plane at right
- Introduce new feature, $z = x^2 + y^2$ (**radial kernel**)
- Map 2 dimensional data onto 3 dimensional data. Now a hyperplane is easy to find. (Imagine slicing a cone!)



Kernels



SVM has MANY Hyperparameters

SVM

C

The “penalty cost”
for misclassifications
(soft margins)

Gamma

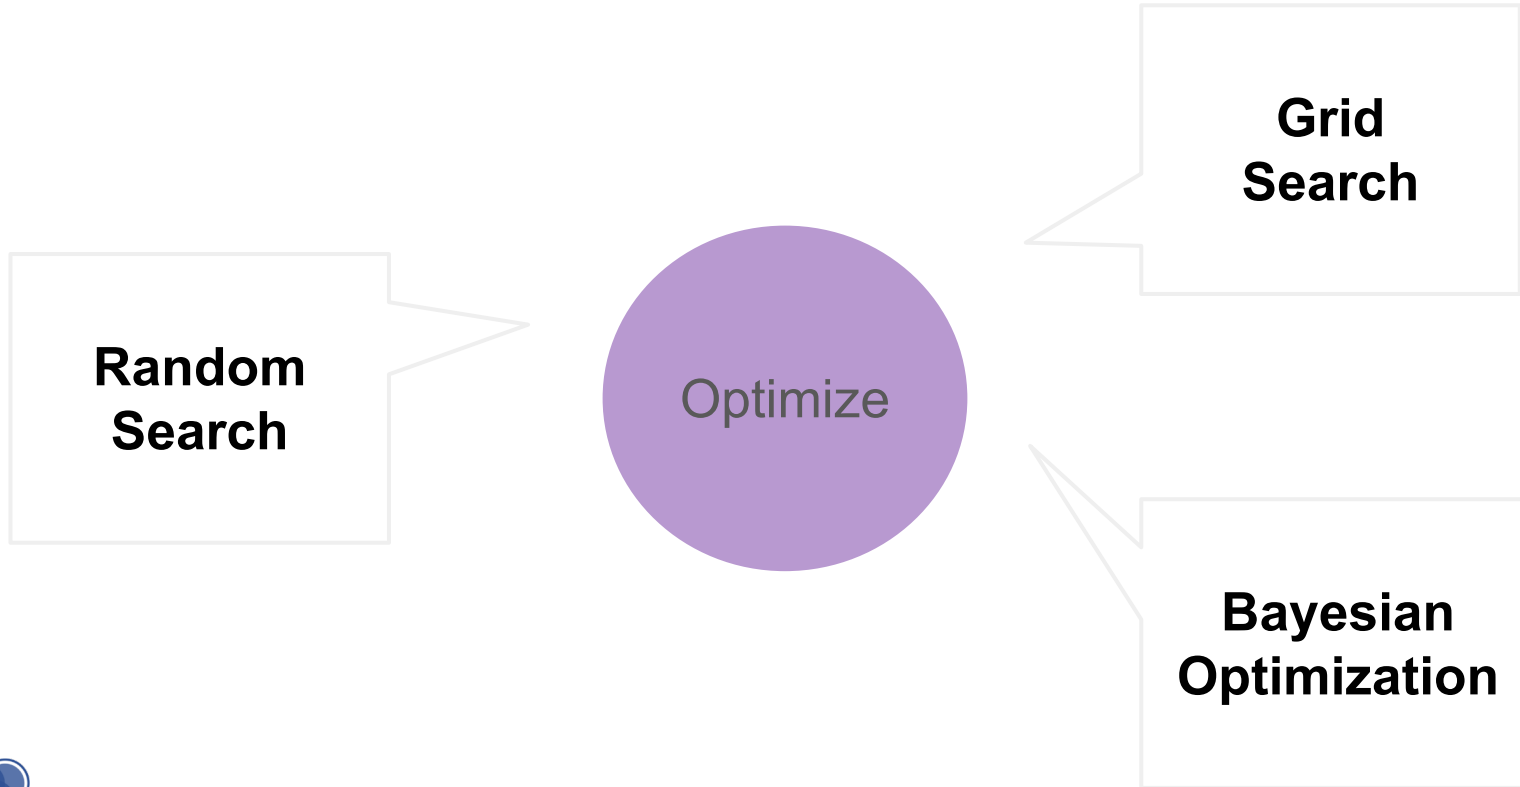
How far the
influence of a single
training example
reaches

Kernels

Method of
transforming our
data set




Finding the Best Hyper Parameters




Curse of Dimensionality

Our search space for the optimal hyper-parameters increases **exponentially** as the number of hyper parameters we are considering increases



Add
dimension



Overview

Perceptron	SVM
<ul style="list-style-type: none">• A very simple model• Will perform poorly if data is not linearly separable	<ul style="list-style-type: none">• More complex model because we have to choose the “penalty cost” associated with misclassifications• Can transform feature space by choosing a Kernel



Demo



More Validation Techniques



Leave-P-out Cross Validation (LpO CV)

Let \mathbf{D} be our whole dataset

Choose a \mathbf{P}

For every combination of \mathbf{P} points in \mathbf{D} :

Use a train/test split with those \mathbf{P} points as test, the rest as train



Leave-P-out: different from K-fold!

Let's say **D** has a size of 4. There are four data points: a , b , c , and d .

K-fold:

- $K = 2$.
- Each fold has a size of 2: $\{a, b\}$ and $\{c, d\}$
- So, we only have 2 possible test sets:
 $\{a, b\}$ and $\{c, d\}$

Leave-P-out:

- $P = 2$.
- We have 6 possible test sets:
 $\{a, b\}$, $\{a, c\}$, $\{a, d\}$, $\{b, c\}$, $\{b, d\}$, and $\{c, d\}$



Leave-P-out Cross Validation (LpO CV)

Pros:

- more fine-grained estimate than k-fold
- test the model's generalization ability

Cons:

- Slow!
 - Runtime increases with larger datasets
 - Runtime explodes with larger P



Monte Carlo Cross Validation

- Need to use **new, random** train/test split each time
 - If you use the same train/test split each time, you're not getting any new information!
- Pros:
 - easy to implement
 - easy to make faster/slower by changing number of iterations
- Cons:
 - random -> train/test splits not guaranteed to be representative of dataset
 - harder to calculate how many iterations you need



The Bootstrap

What if we don't have enough data?

- Use **bootstrap datasets** to approximate the test error
- **Sample with replacement** from the original training dataset (with n samples) to generate **bootstrap datasets** of size n
 - Some data points may appear more than once in the generated data
 - Some data points may not appear
- Estimate of test error = average error among bootstrap datasets



Demo



Coming Up

- **Assignment 8:** Due midnight of next class
- **Next Lecture:** Applications of *Unsupervised* Learning



CDS Education

We explore, learn, and educate big minds.