

# INFO 1998: Introduction to Machine Learning

Pull up Lecture 5 Demo from website as well!



**CDS Education**

We explore, learn, and educate big minds.

# Lecture 5: Fundamentals of Machine Learning Pt. 2

INFO 1998: Introduction to Machine Learning

## Tuning Models



**CDS Education**

We explore, learn, and educate big minds.

# Announcements

## Mid-Semester Check-in

Where you should be right now:

- Have an idea of what your problem statement/hypothesis is
- Have your group chosen
- Have your data set chosen and some progress

Drop Deadline: **March 20<sup>th</sup>**



# What We'll Cover

**Last Time's Goal:** identify what ML is and write ML code (to some extent)

**This Time's Goal:** how to tell if your ML model is *useful*



# Agenda

1. Review
2. Measuring Accuracy
3. Bias-Variance trade-off
4. Feature Selection
5. Other Types of machine learning




# Review: Defining ML

We want to predict the future

- Take some known input and output
- Learn the data's pattern and come up with a way to, given a future input, predict the corresponding output



# Review: Model



- Takes in input and output and learns the relationship
- Used to predict outputs
- “Model training” = learn a relationship/program
- “Model validation” = see if the learned relationship is accurate on other data
- “Model testing” = final model performance



# Measuring Bias / Loss


(training accuracy)






# Loss, Cost, and Score Functions

**Loss Function:** Penalty for missing a single data point



# Loss, Cost, and Score Functions

**Cost Function:** Indicates how bad the whole model is



- Applies loss function to each point, then combines that into a single number
  - ex: average of (loss from each point)



# Loss, Cost, and Score Functions

**Cost Function:** Indicates how bad the whole model is



- Applies loss function to each point, then combines that into a single number
  - ex: average of (loss from each point)



# Loss, Cost, and Score Functions

**Cost Function:** Indicates how bad the whole model is



- Applies loss function to each point, then combines that into a single number
  - ex: average of (loss from each point)



# Loss, Cost, and Score Functions

**Cost Function:** Indicates how bad the whole model is



- Applies loss function to each point, then combines that into a single number
  - Ex:, Total Loss, Average Loss



# Loss, Cost, and Score Functions

- **Score Function**

- A more interpretable version of the cost function (how well we did)
- Loss/Cost used in training to help a model learn, Score is just what we use for interpretability

Cost =   Score = 50% Accuracy



# Linear Regression Loss Formula: Euclidean Distance


$$\text{loss}(x, y^*) = (h(x) - y^*)^2$$

Two things to note about this loss function:

- Positives and negatives won't cancel
- Large errors are penalized to a power of 2 more

**Cost Function** - average of the loss function over all the points

LSS Model Gives Residual Sum of Squares (RSS)




# Linear Regression Loss Formula: Euclidean Distance

$$\text{loss}(x, y^*) = (h(x) - y^*)^2$$

In what situations might you want a high penalty loss function as opposed to a lower penalty function?

- High stakes situations (Ex: Cancer Diagnosis)
- Data does not have many outliers

LSS Model Gives Residual Sum of Squares (RSS)





## Solution: Compare to Baseline

- When determining accuracy, usually want to compare our model to a **baseline**
  - For regression, one baseline model is the model that predicts the **average** of the target value for every point
  - For our purposes: don't worry about the baseline *model*, just have a set of baseline *predictions*



# Cost to Accuracy Score

- sklearn's score function is:

$$1 - ([\text{Cost of model}] / [\text{Cost of baseline}])$$

- 1 is very, very good
- 0 means you were as bad as the baseline
- <0 means either your baseline predictions were accurate, or you really, really messed up
- **NOT USED IN TRAINING, For our interpretability**




## Question!


- Let's say we have a dataset  $\{(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)\}$ .
- Suppose our model outputs the following:
  - $y_i$ , if  $x = x_i$  for some  $i$  in  $\{1, 2, \dots, n\}$
  - $-3.14159$  otherwise
- What is the cost of this model (using Euclidean distance)? Hint: the answer is quite simple.
- Should we expect this algorithm to perform well in predicting outputs for new inputs?



# Training Data



## Cost = 0, but model is horrible...




MORAL: Assumptions are important!

# No Free Lunch Theorem

**Every ML algorithm makes assumptions!**

Ex: Linear regression assumes data has a linear relationship



MORAL: Assumptions are important!



# Overfitting and Underfitting

(what makes a model good?)



# Model Goals

When training a model, we want our model to:

- Capture the trends of the training data
- Generalize well to other samples of the population
- Be moderately interpretable

The first two are especially difficult to do simultaneously!

The more sensitive the model, the less generalizable and vice versa.





# Model Goals

When training a model, we want our model to:


- Capture the trends of the training data
- Generalize well to other samples of the population
- Be moderately interpretable

The first two are especially difficult to do simultaneously!


The more sensitive the model, the less generalizable and vice versa.




**Underfitting**: A situation when your model is **too simple** for your data.




# Underfitting




# Underfitting




# Underfitting




# Underfitting



# Underfitting




# Underfitting







# Underfitting: at least the models are consistent...




**Overfitting**: A situation when your model is **too complex** for your data.




# Overfitting




# Overfitting




# Overfitting




# Overfitting



# Overfitting




# Overfitting: Inconsistent Models!






# Overfitting: Results from training with high sensitivity



# Overfitting: doesn't generalize well!



# Bias and Variance



# Definitions

## Bias

- A measure of underfitting

## Variance

- A measure of overfitting

Either alone is hard to interpret, but together they are helpful

<http://www.r2d3.us/visual-intro-to-machine-learning-part-2/>



# Balancing Bias and Variance

$$\mathbb{E}\left[(y - \hat{f}(x))^2\right] = \text{Bias}[\hat{f}(x)]^2 + \text{Var}[\hat{f}(x)] + \sigma^2$$

$$\text{Bias}[\hat{f}(x)] = \mathbb{E}[\hat{f}(x) - f(x)]$$

$$\text{Var}[\hat{f}(x)] = \mathbb{E}[\hat{f}(x)^2] - \mathbb{E}[\hat{f}(x)]^2$$



Error = (expected loss of accuracy)<sup>2</sup> + inconsistency of model + irreducible error



Low Variance

High Variance



Low Bias



Overfitting

Underfitting

High Bias



## What does this mean intuitively?

### Bias

- Bad
- Results from incorrect assumptions in the learning algorithm

### Variance

- Bad
- Results from sensitivity to fluctuations in the data



# What can you do to reduce Bias and Variance?

## Bias

- Increase model size
- Change type of model
- Add new features


## Variance

- Add more data
- Decrease model size
- Reduce features





# Balancing Bias and Variance



# Feature Selection

(adjusting models)



# Methods

- **Goal:** Find subset of features that gives a good enough model, in a reasonable amount of time.



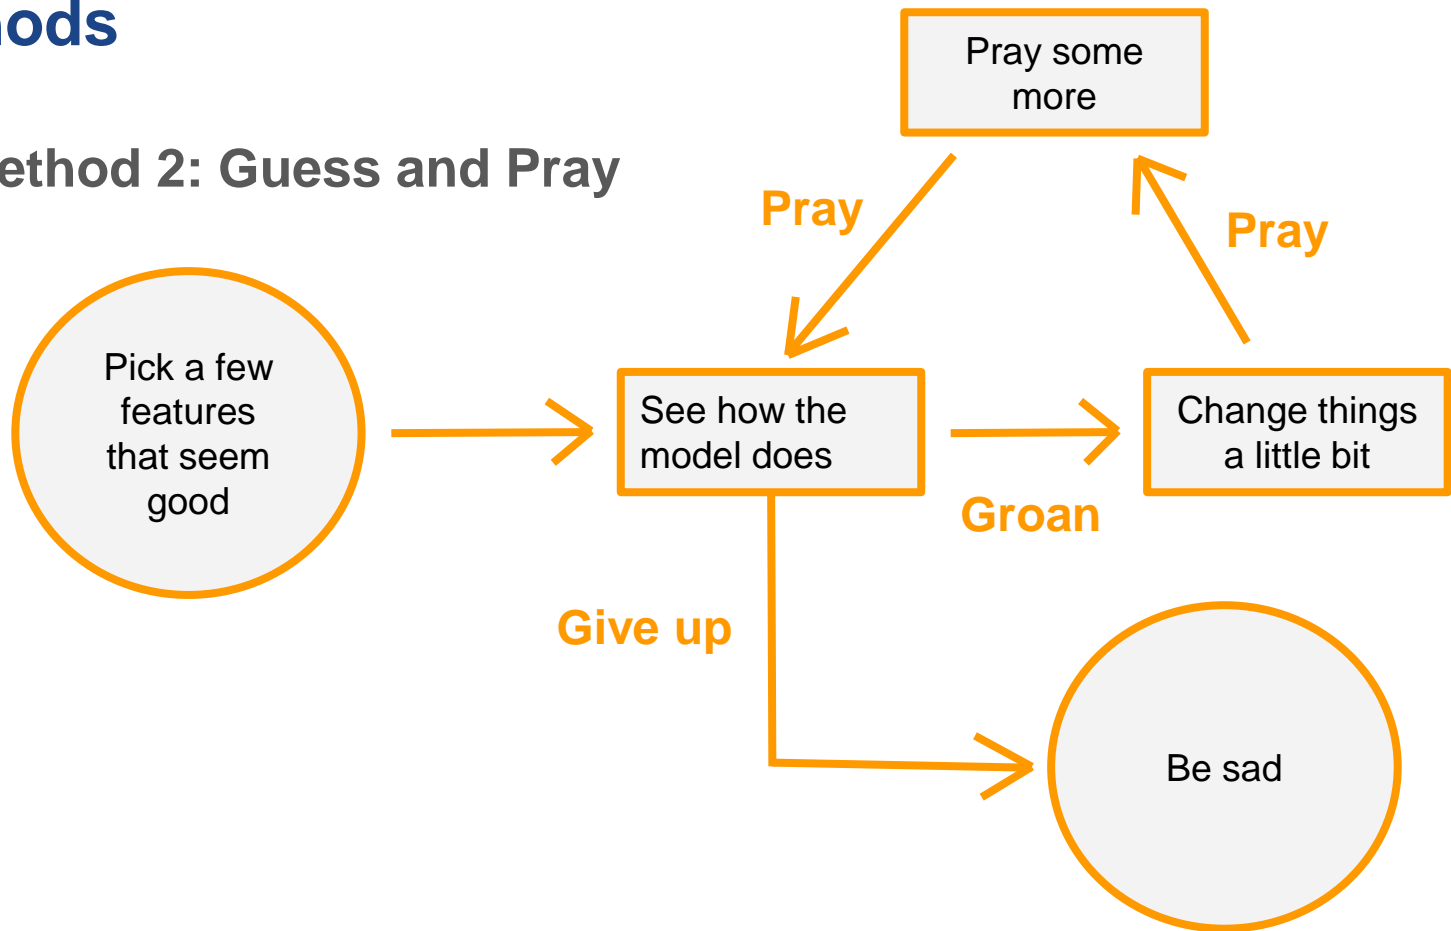
# Methods

- **Goal:** Find subset of features that gives a good enough model, in a reasonable amount of time.
- **Method 1: Best Subset**
  - Test **all** subsets for best one
  - Benefits:
    - **Best** subset out of current features
  - Drawbacks:
    - Slow
    - Even slower with feature engineering



# Methods

- Method 2: Guess and Pray



# Methods

- **Goal:** Find subset of features that gives a good enough model, in a reasonable amount of time.
- **Method 2: Guess and Pray**
  - Guess
  - Benefits:
    - ??
  - Drawbacks:
    - Time consuming for data scientist
    - Unreliable



# Methods


- **Goal:** Find subset of features that gives a good enough model, in a reasonable amount of time.
- **Method 3: Stepwise**
  - Pick a few features, then programmatically add/remove features using statistics
  - Benefits:
    - Complexity and runtime are adjustable
  - Drawbacks:
    - Can do very badly if you're not careful
    - Requires more thinking



# Correlation, $r$

The correlation between two variables describes to what extent changing one would change the other.

- Real-valued in  $[-1,1]$
- A variable is always perfectly correlated with itself (correlation=1)





## Important Case: Collinearity

**Collinear:** when two features have a correlation near -1 or 1

- If a feature is collinear with the target, then it's a good choice for linear regression
- If two features are collinear, they're *redundant*
  - Might as well not use one of them
  - Some models *require/assume* no collinear features
  - Takes more time, and doesn't add much information at the cost of *increased variance/sensitivity*



## Side Note: Scaling and Normalizing


- Some models require data to be centered
- Some models need features to be on the same scale
  - Divide by max, minus min divide by max minus min, minus mean divide by standard deviation.



## Side Note: Scaling and Normalizing

- Some models require data to be centered
- Some models need features to be on the same scale


**Divide by max:** Bounds data  $\leq 1$



## Side Note: Scaling and Normalizing


- Some models require data to be centered
- Some models need features to be on the same scale

-  $\frac{\text{min}}{(\text{max} - \text{min})}$  : Bounds data between [0,1]



## Side Note: Scaling and Normalizing

- Some models require data to be centered
  - Some models need features to be on the same scale
- **mean / standard deviation:** Z scores – Distance from mean




## Other Ways to Adjust your Model

- Hyper Parameters – Ex: Learning rate, etc.
- Feature engineering – Ex: Manipulating dataset
- Just changing to a different algorithm



# Demo



# Different Types of ML

(supervised & unsupervised)  
(classification & regression)






# Supervised vs. Unsupervised

## Supervised learning...

- **Known target variable info**
- **Validation examples**

Classification




Supervised learning

## Unsupervised learning...

- **Unknown target variables**
- **Difficult to validate**


Clustering




Unsupervised learning



# Classification vs. Regression



Classification  
Discrete Class Labels



Regression  
Continuous Quantity



## Other Classes of ML Algorithms (which we won't cover)

- What if you can't / don't want to see all your data at once?
- Maybe you only want to use a few pieces of your data (but don't have the time to manually select each piece of data...)
- A different approach, Trial & Error: The algorithm tries one thing, sees how that works, makes adjustments, tries again, etc.



# Final Notes



*Always remember both bias and variance!*

# Coming Up

- **Assignment 5:** Due at midnight on March 22<sup>nd</sup> , 2023
- **Next Lecture:** Intro to Classification
- Last day to drop: March 20<sup>th</sup> (Next Monday)



**CDS Education**

We explore, learn, and educate big minds.