INFO 1998: Introduction to Machine Learning

CDS Education

We explore, learn, and educate big minds.

Mid-Semester Feedback

- How are we doing? Improvements/suggestions are welcome!
- Google form will be posted on Ed

Mid-Semester Check-in for final project

- between lecture 6 and 8
- Need Dataset + Jupyter Notebook (basic data cleaning/exploration)
- 5 points of the final grade
- (Full announcement on Ed)

Lecture 6: Intro to Classifiers

INFO 1998: Introduction to Machine Learning

CDS Education

We explore, learn, and educate big minds.

Agenda

1. What is a Classifier?
2. K-Nearest Neighbors Classifier
3. Fit/Overfitting
4. Confusion Matrices

What are Classifiers?

What are Classifiers?

Classifiers are able to help answer questions like…

- “What species is this?”
- “What major is a student in based on their classes?”
- “Which Hogwarts House do I belong to?”
- “Am I going to pass this class?”

What are Classifiers?

- Classifiers predict the class/category of a set of data points. This class/category is based off of the target variable we are looking at.

- Difference between linear regression and classifiers
  - Linear regression is used to predict the value of a **continuous variable**
  - Classifiers are used to predict **categorical or binary variables**

What are Classifiers?

Two categories of classifiers: lazy learners and eager learners

- **Lazy Learners**
  - Store the training data and wait until a testing data appear
  - Classification is conducted based on the most related data in the stored training data
  - Less training time, more time in predicting
- **Eager Learners**
  - Construct a classification model based on the given training data before receiving data for classification
  - More training time, less time in predicting

K-Nearest Neighbors Classifier

What is the KNN Classifier?

- Lazy learner classifier
- Easy to interpret
- Fast to calculate
- Good for coarse analysis

How Does It Work?

Uses the k (a user specified value) nearest data points to predict the unknown one

- A simple assumption: the values **nearest** to a data point are **similar** to it
- k is a **hyperparameter** of the KNN model (a parameter that affects the learning process)!

How Does It Work?

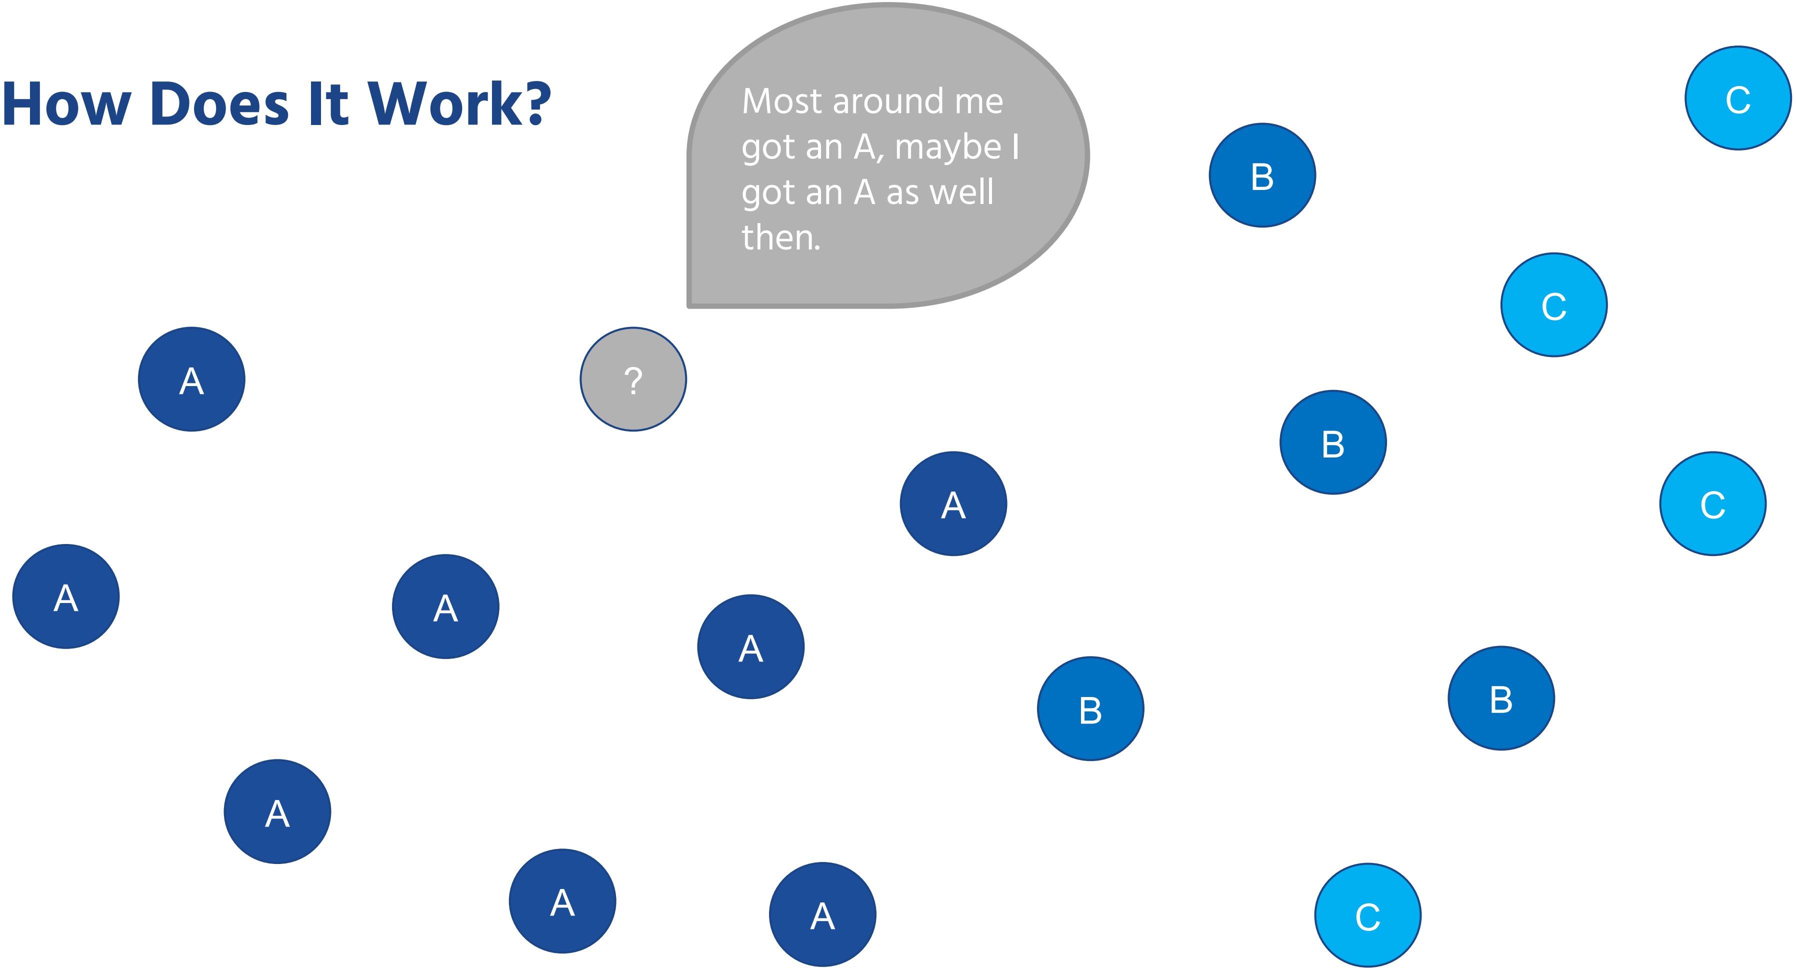

How Does It Work? (Step-By-Step Example)

Define a *k* value (in this case k = 3)

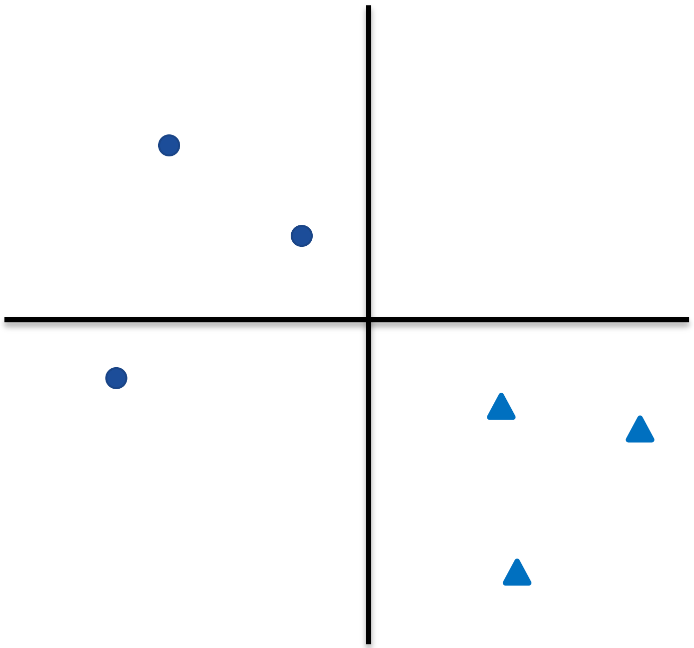

How Does It Work? (Step-By-Step Example)

Define a *k* value (in this case k = 3)

Pick a point to predict (blue diamond)

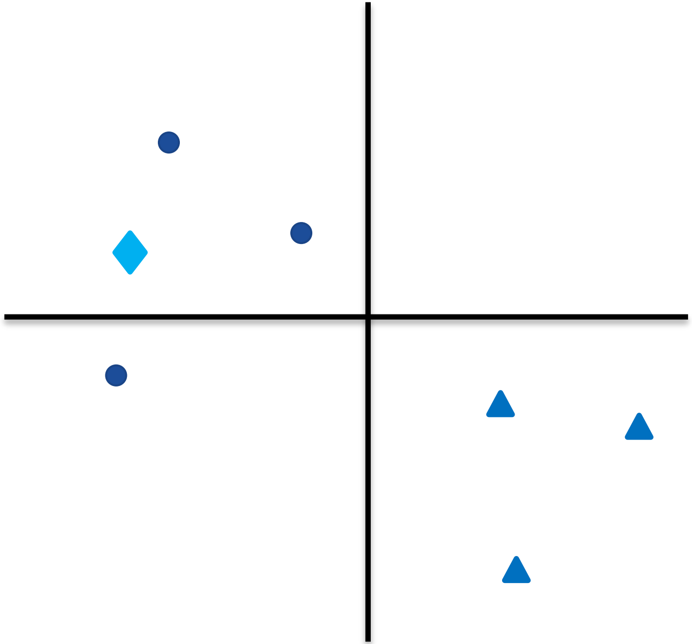

How Does It Work? (Step-By-Step Example)

Define a *k* value (in this case k = 3)

Pick a point to predict (blue diamond)

Count the number of closest points

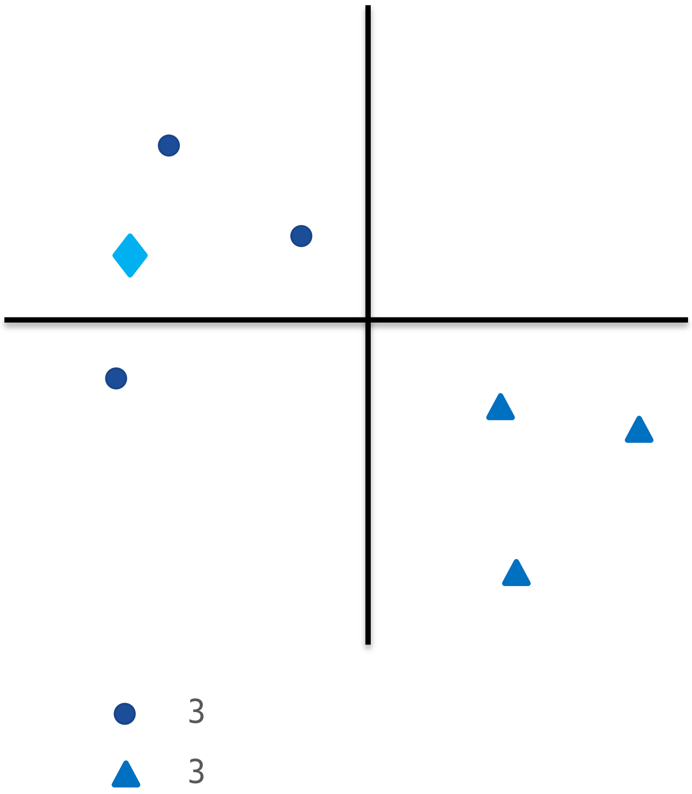

How Does It Work? (Step-By-Step Example)

Define a *k* value (in this case k = 3)

Pick a point to predict (blue diamond)

Count the number of closest points

Increase the radius until the number of points within the radius adds up to 3

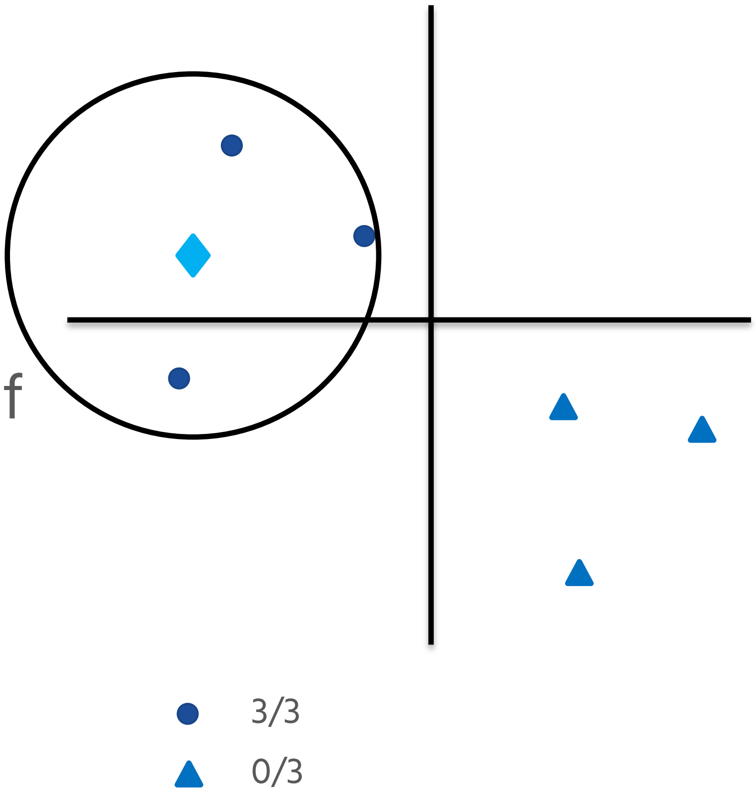

How Does It Work? (Step-By-Step Example)

Define a *k* value (in this case k = 3)

Pick a point to predict (blue diamond)

Count the number of closest points

Increase the radius until the number of points within the radius adds up to 3

Predict the blue diamond to be a blue circle!

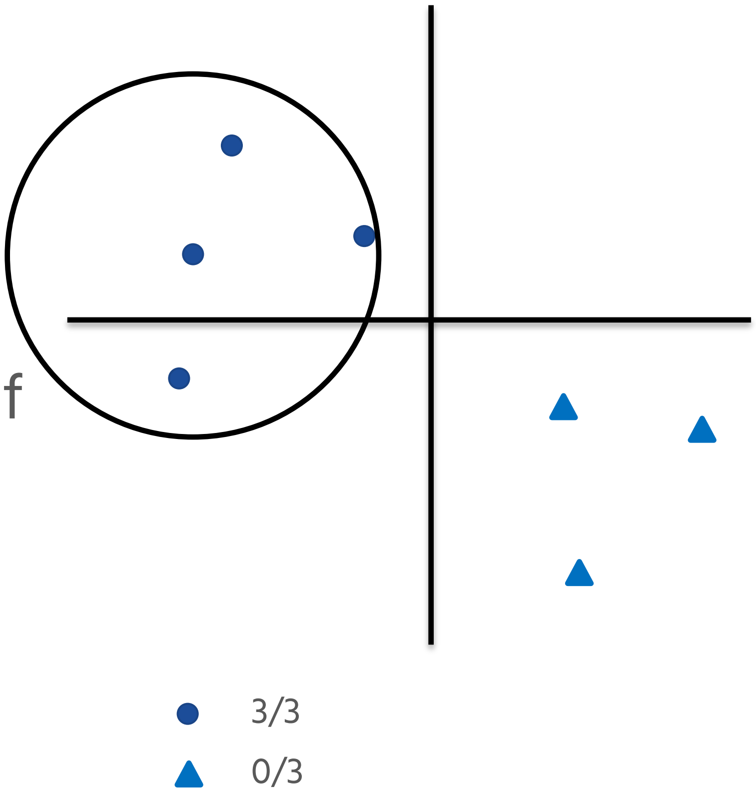

Demo

Fit/Overfitting

Overfitting

When the model corresponds too closely to training data and then isn't transferable to other data.

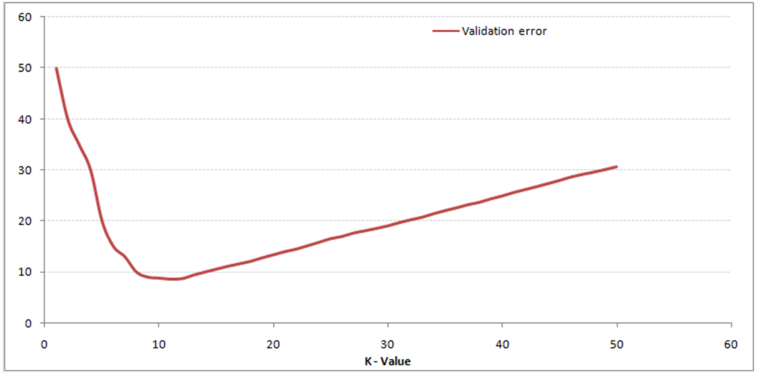

Can fix by:

- Splitting data into training and validation sets
- Increasing k

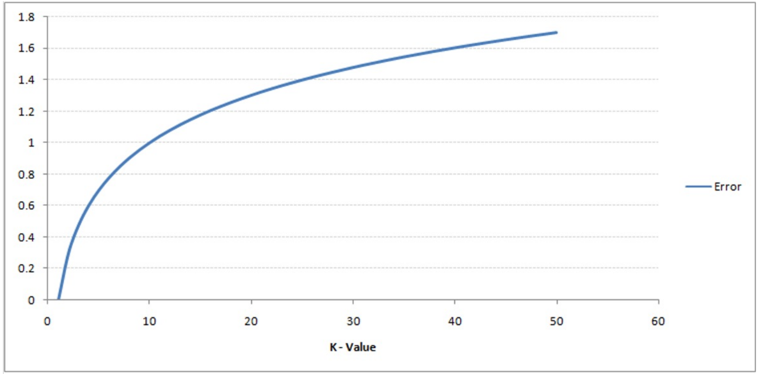

Relationship Between k and Fit

The **k** value you use has a relationship to the fit of the model

A higher k gives a smoother line, but too large of a k and it is the average of all the data (or the label that is most common/likely)

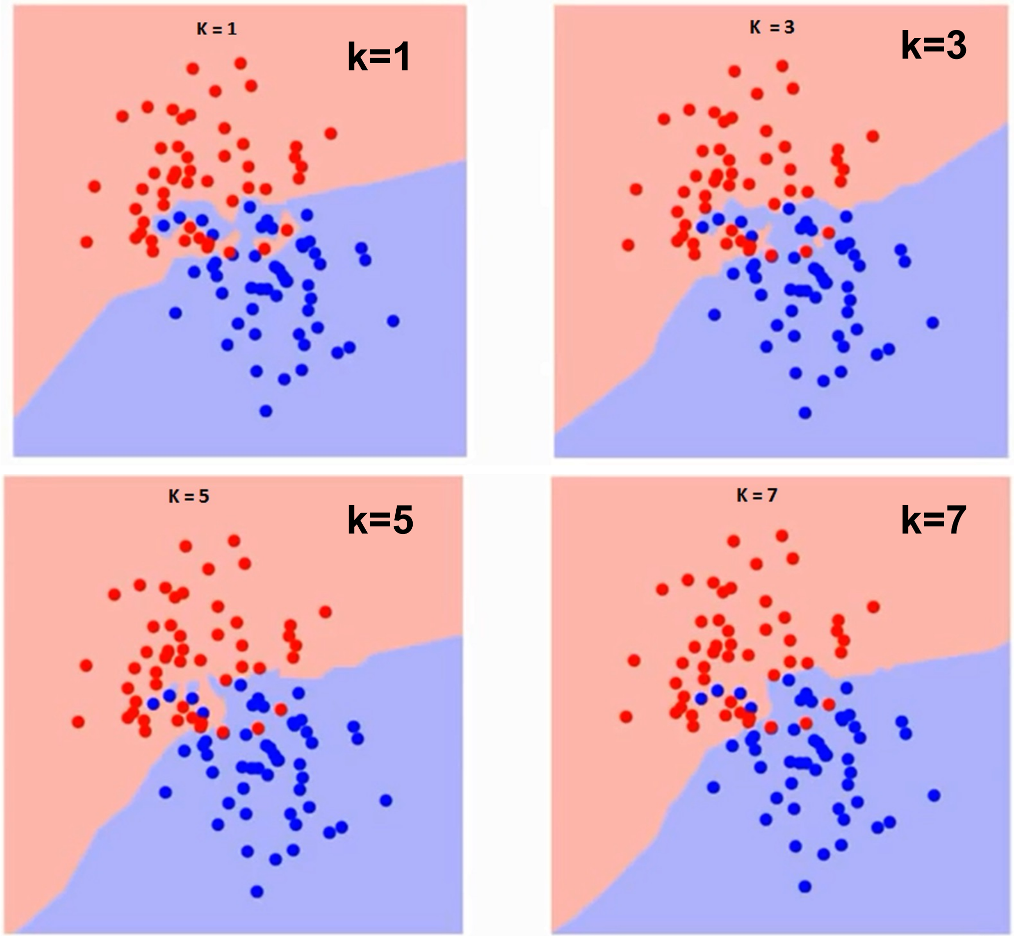

Confusion Matrix

What is a Confusion Matrix?

Table used to describe the performance of a classifier on a set of binary test data for which the true values are known

	p' (Predicted)	n' (Predicted)
p (Actual)	True Positive	False Negative
n (Actual)	False Positive	True Negative

Sensitivity

Called the **true positive rate**

Tells us how many positives are correctly identified as positives

Optimize for: Initial diagnosis of fatal disease

	p' (Predicted)	n' (Predicted)
p (Actual)	True Positive	False Negative
n (Actual)	False Positive	True Negative

Sensitivity = True Positive/ (True Positive + False Negative)

Specificity

Called the **true negative rate**

Tells us how many negatives are correctly identified as negatives

Optimize for: testing for a disease with a risky treatment

	p' (Predicted)	n' (Predicted)
p (Actual)	True Positive	False Negative
n (Actual)	False Positive	True Negative

Specificity = True Negative/ (True Negative + False Positive)

Question [pollev.com/vchen100]

Which is an example of when you would want **higher specificity**?

A. DNA tests for a death penalty case
B. Testing if dining hall food is safe and clean
C. Airport security
D. Testing if someone is positive (covid)

Overall Accuracy

Proportion of correct predictions

	p' (Predicted)	n' (Predicted)
p (Actual)	True Positive	False Negative
n (Actual)	False Positive	True Negative

Accuracy =
(True Positive + True Negative) / Total

Overall Error Rate

Proportion of incorrect predictions

	p' (Predicted)	n' (Predicted)
p (Actual)	True Positive	False Negative
n (Actual)	False Positive	True Negative

Error = (False Positive + False Negative) / Total

Precision

Proportion of correct positive predictions among all positive predictions

	p' (Predicted)	n' (Predicted)
p (Actual)	True Positive	False Negative
n (Actual)	False Positive	True Negative

Precision = True Positive / (True Positive + False Positive)

Coming Up

- **Assignment 6**: Due on Oct 26th, 2022
- **Mid-Semester Check-In:** More details on ED Discussion!
- **Feedback Survey:** Please fill it out!
- **Web-scraping workshop:** 10/25 4:30-5:30pm Hollister 368
- **Next Lecture**: Applications of Supervised Learning

CDS Education

We explore, learn, and educate big minds.