Lecture 5: Fundamentals of Machine Learning Pt. 2

INFO 1998: Introduction to Machine Learning

Bias vs. Variance & Tuning Models

CDS Education

Announcements

Mid-Semester Check-in

Where you should be right now:

- Have an idea of what your problem statement/hypothesis is
- Have your group chosen
- Have your data set chosen and some progress

Complete in OH or after lecture anytime between **now** and **Oct 23rd**

Cornell Drop Deadline: **Oct 21**

Apply to Cornell Data Science! 📣

- All subteams are recruiting freshmen this semester!
  - Deadline: **October 17th, 11:59pm**
  - Don't forget to also submit the College of Engineering application.

- Application Link: https://cornelldata.science/recruitment

- If you're enjoying this class…
  - you'll LOVE being on CDS 👀

Subteam UTea trip!

What We'll Cover

Last Time's Goal: identify what ML is and write ML code (to some extent)

This Time's Goal: how to tell if your ML model is *useful (good)*

Agenda

1. Review
    - Types of Machine Learning
2. Measuring Accuracy/Error
3. Model Selection
4. Feature Selection

Review: Defining ML

We want to predict the future

- Take some known input and output
- Learn that data's "pattern" to:
  - Given a future input, predict[1] the corresponding output

[1] We model how the output is generated

Review: Machine Learning Pipeline

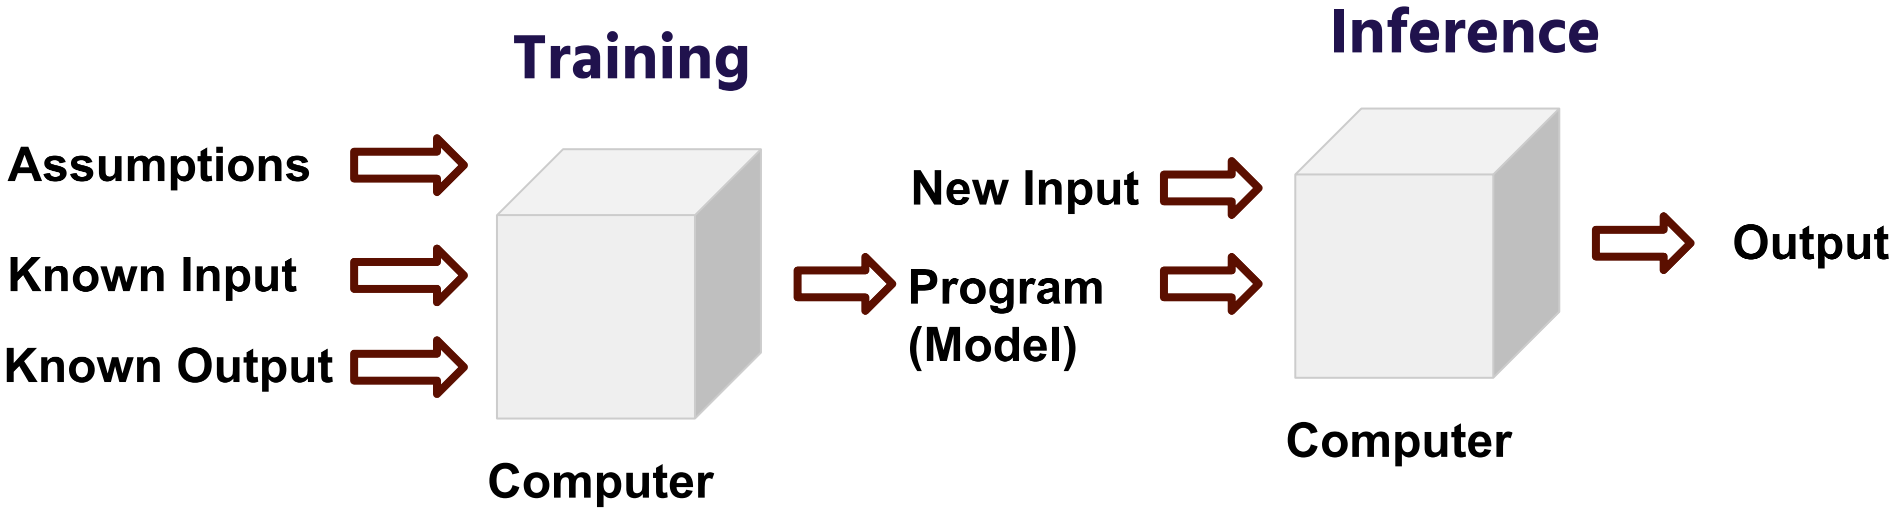

Review: Model

- “Model training” = learn a relationship
- “Model testing” = check if the learned relationship is generalizes
- “Model validation” = simulates model performance when used in real life

Different Types of ML

(supervised & unsupervised)
(classification & regression)

Supervised vs. Unsupervised

Supervised learning…

- Goal: Predict output
- Needs known output/target

Unsupervised learning…

- Goal: learn more about the data (ex. trends)
- Doesn't need known output

Examples of Supervised: Classification and Regression

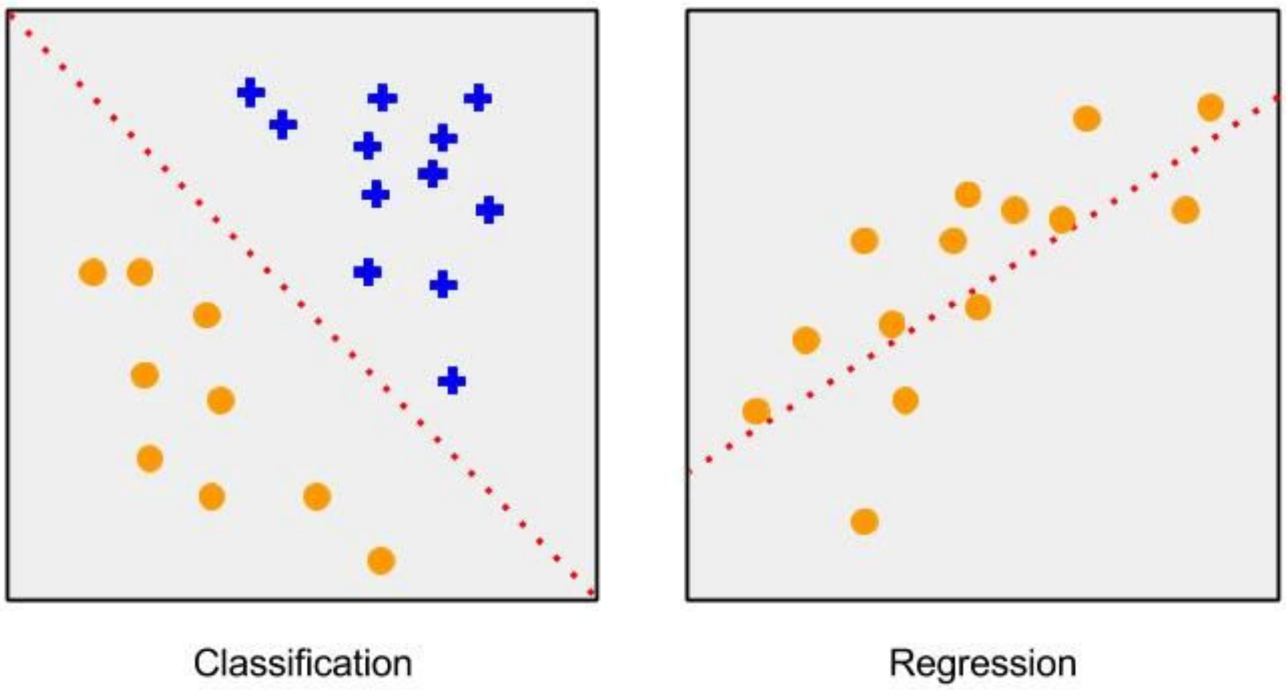

Classification or Regression?

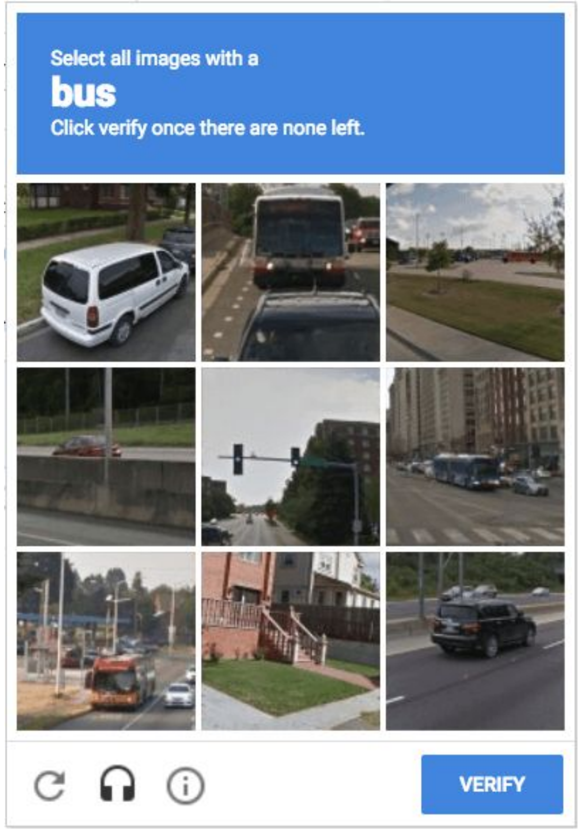

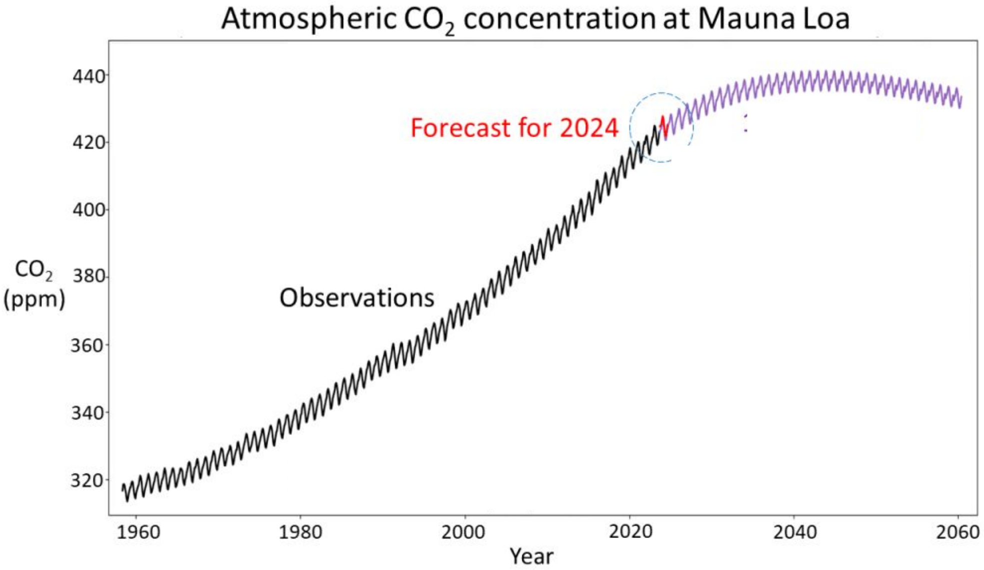

Classification or Regression? Examples from my internship

Detecting fake students
(adults using student discount)

Predicting the value of a customer

Measuring Training Accuracy

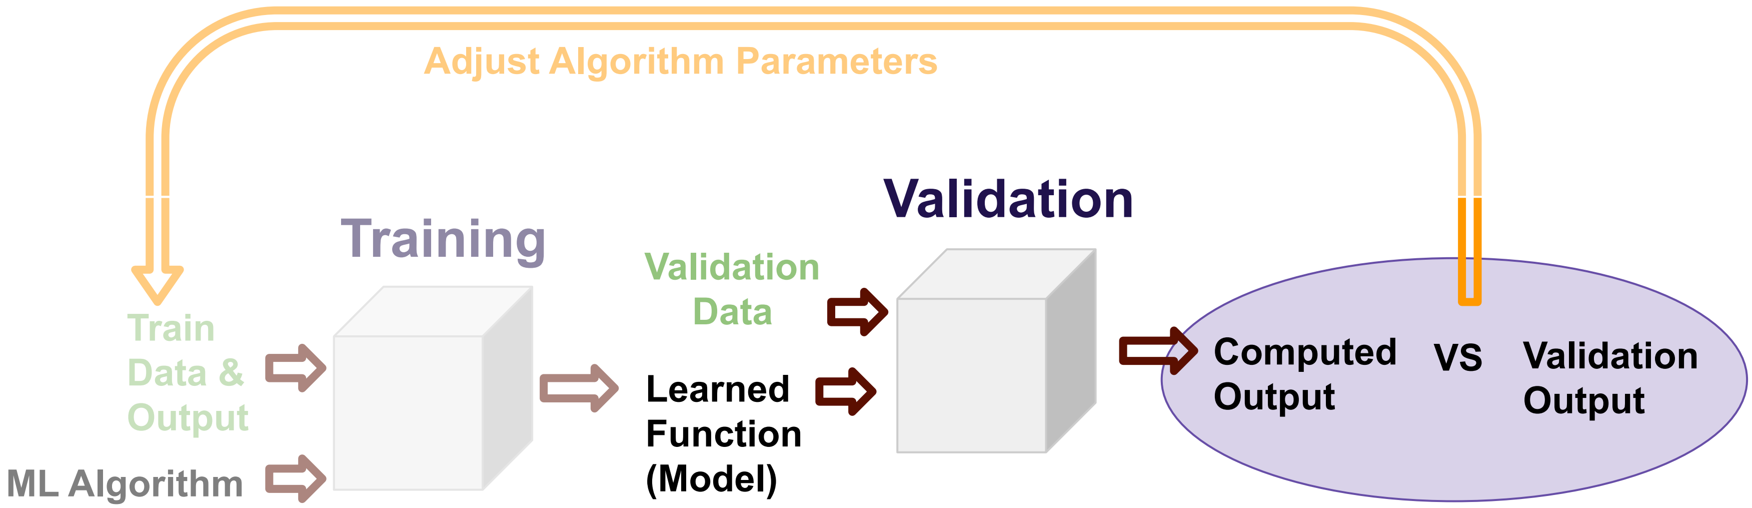

1. Split data (lecture 7)
2. Assess model accuracy (today)
3. Adjust Model (a bit today)

Loss, Cost, and Score Functions

- **Loss Function**
  - How far is a prediction from its corresponding answer
  - Used as a penalty for mislabelling in training to help a model learn

- **Cost**
  - Applies loss function to each point, then combines that into a single number

- **Metric (Score Function)**
  - How well the model did across all data points
  - Interpretable, for the model builder

Examples of Loss & Metrics: Multiple Choice Exams

- How would you evaluate these?
  - If the answer is A) but you pick B)

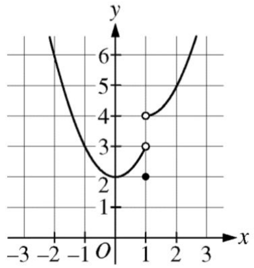

Graph of f

. The graph of the function f is shown in the figure above. The value of $\lim_{x \to 0} f\left(1 - x^2\right)$ is

(A) 1 (B) 2

9

Why does Akira say his meeting with Chie is "a matter of urgency" (line 32)?

A) He fears that his own parents will disapprove of Naomi.

B) He worries that Naomi will reject him and marry someone else.

10

Which choice provides the best evidence for the answer to the previous question?

A) Line 39 ("I don't . . . you")

B) Lines 39-42 ("Normally . . . community")

Examples of Loss & Metrics: Multiple Choice Exams

- Zero-one loss:
    - 1 if prediction == answer
    - 0 if prediction == answer

Examples of Loss & Metrics: Google Maps

- How would you evaluate this?
  - If Google Maps says it will take 26 mins but it actually takes x minutes

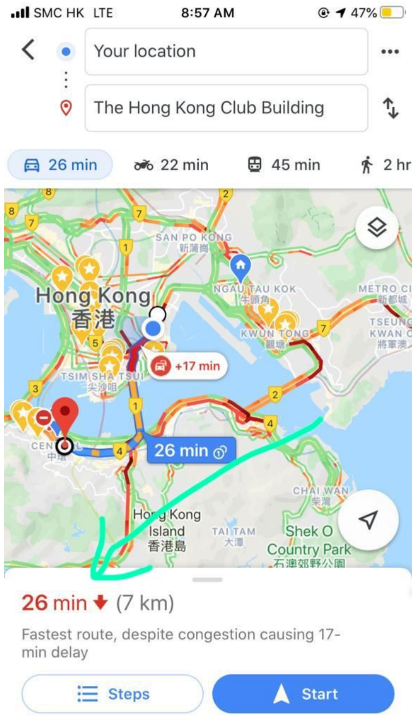

Linear Regression Loss Formula: Euclidean Distance

$$\text{loss}(x_i, y_i) = (h(x_i) - y_i)^2$$

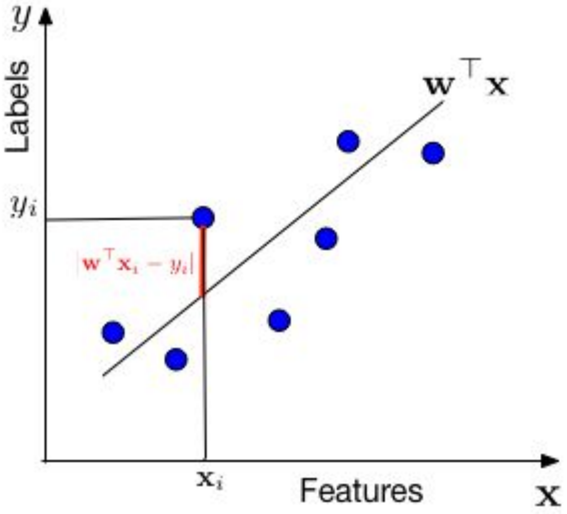

Two things to note about this loss function:

- Positives and negatives won't cancel
- Large errors are penalized to a power of 2 (more)

In what situations might you want a low penalty loss function as opposed to this high penalty loss function?

Linear Regression Loss Formula: Euclidean Distance

$$\text{loss}(x_i, y_i) = (h(x_i) - y_i)^2$$

What could the **cost function** be?

- MSE = (...)/N
  - Where N is the number of data points

How do you know if something is good?

- “I throw at a speed of 35 ft/sec.”

How do you know if something is good?

- “I throw at a speed of 35 ft/sec. The average for pros is 27 ft/sec.”

Compare to Baseline

- When evaluating accuracy, compare our model to a **baseline**
  - For regression, one baseline model is the model that predicts the **average** of the target value for every point
  - For our purposes: don't worry about the baseline *model*

Sk-learn's score function

1 - ([Cost of model] / [Cost of baseline])

- **>0** means you beat the baseline
- **0** means you were equal to the baseline
- **<0** means you're worse than the baseline

Overfitting and Underfitting

(how generalizable is the performance?)

Model Goals

When training a model, we want our model to:

- Capture the trends of the training data sample
- Generalize well to the whole population
- Be moderately interpretable

The first two are especially difficult to do simultaneously!

- Want to choose the right amount of complexity

Generate Samples To Illustrate Over/Under fitting

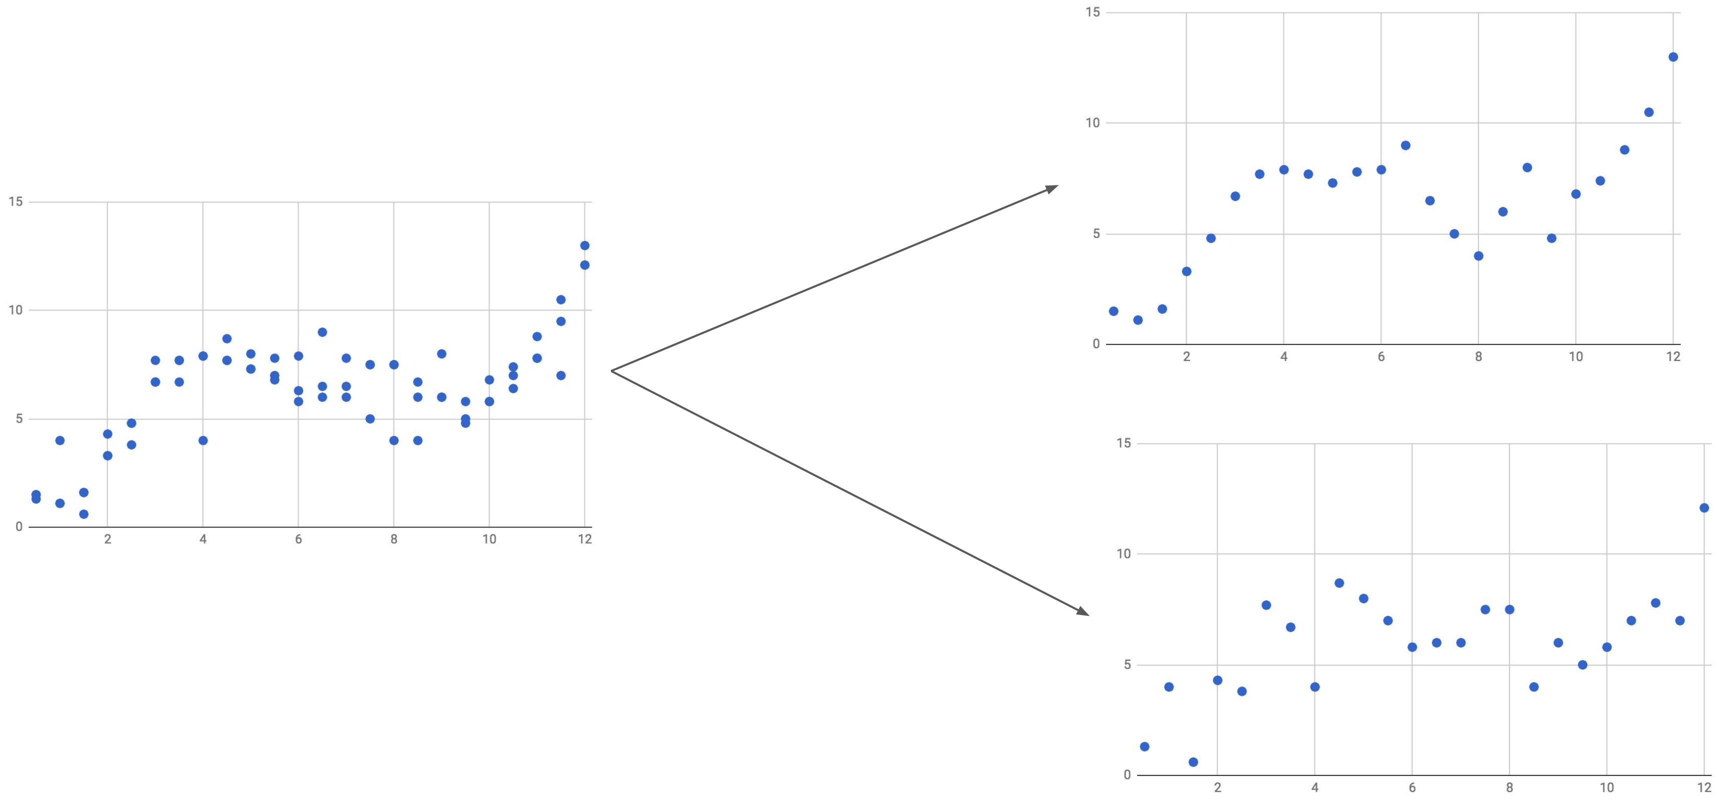

Underfitting

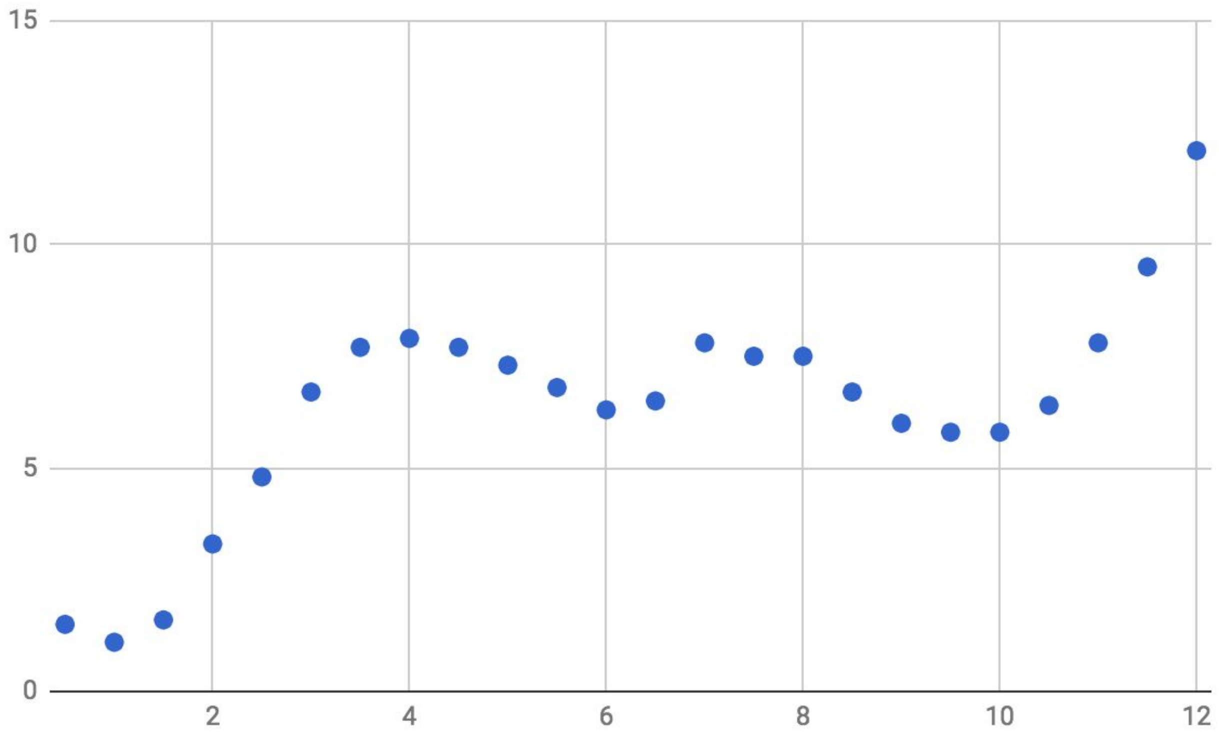

Underfitting: Too simple

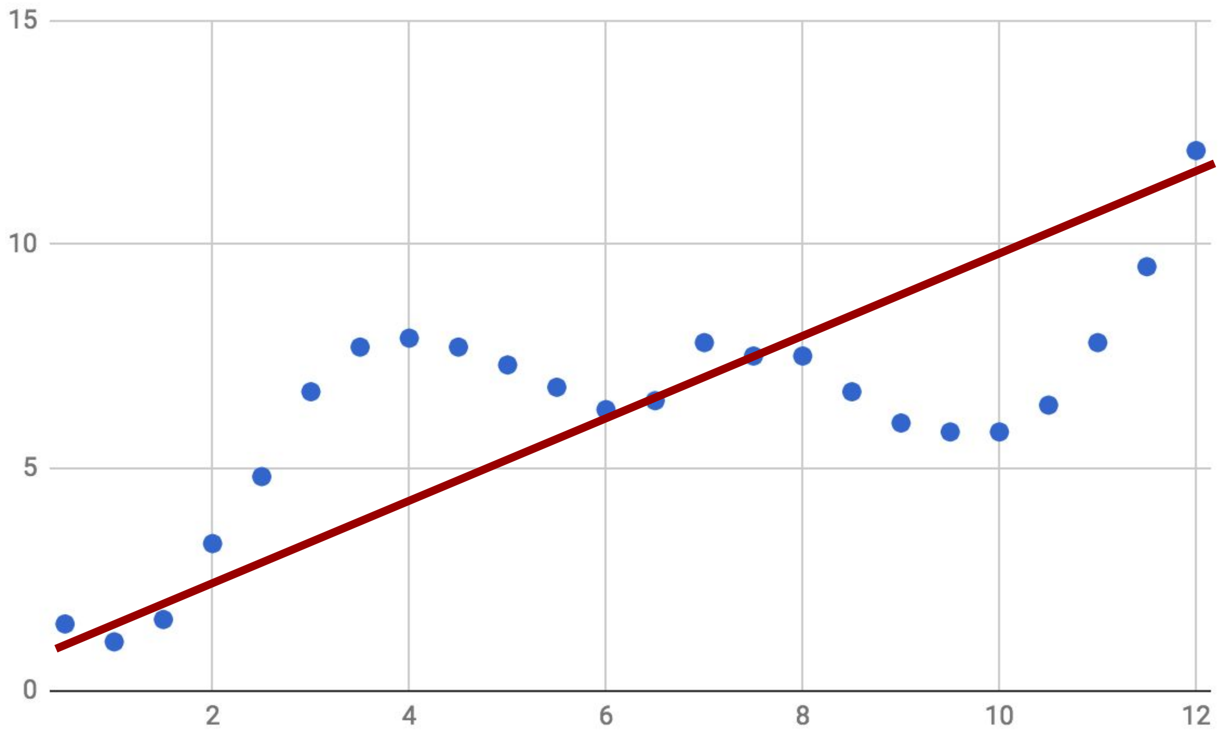

Underfitting

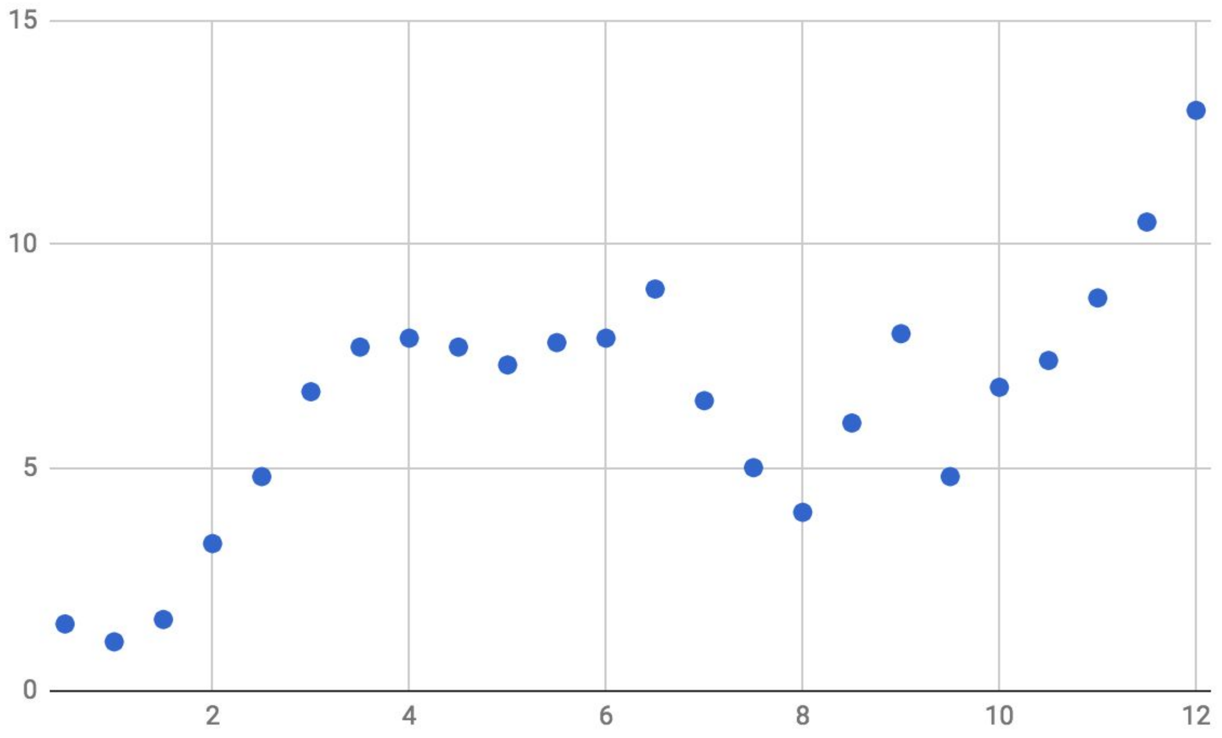

Underfitting: Too simple

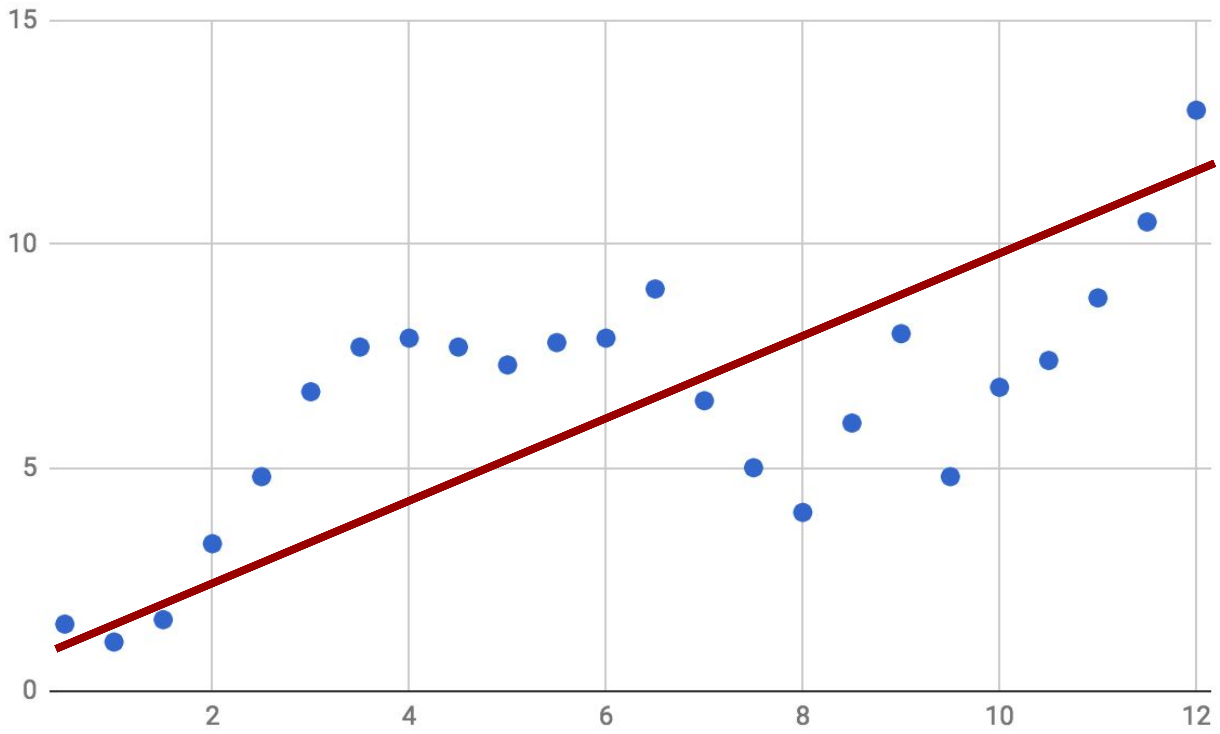

Underfitting

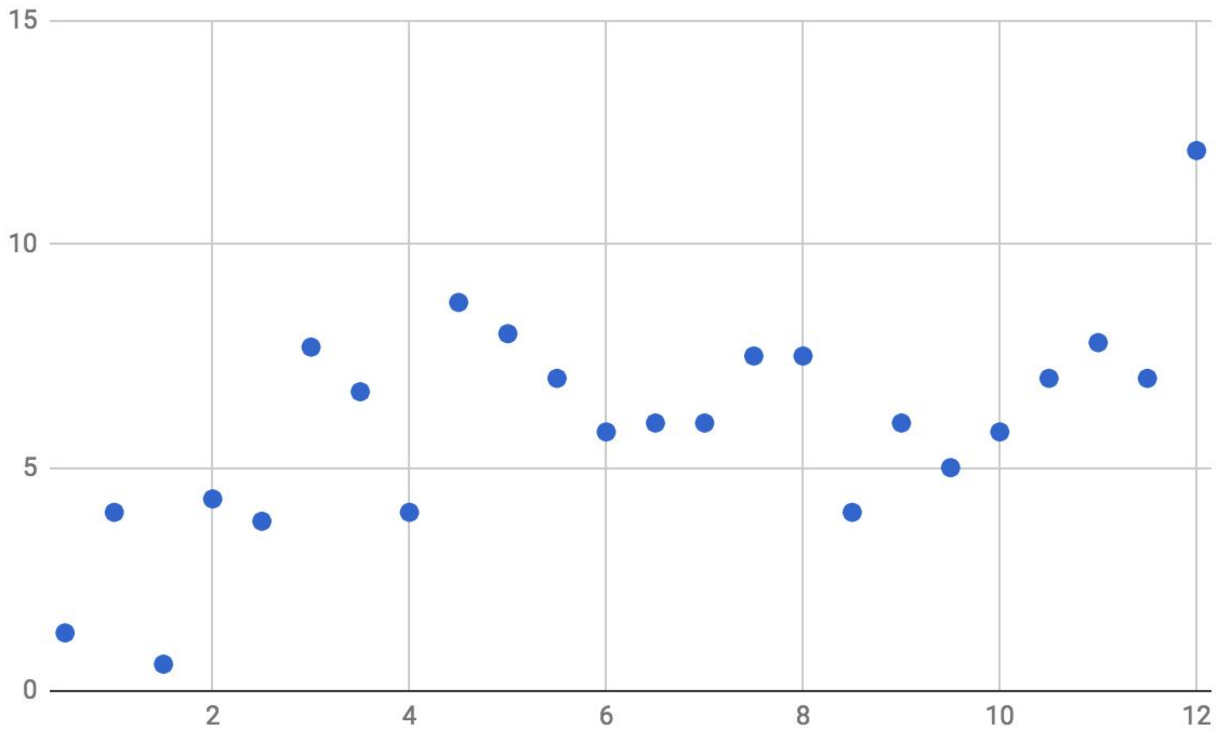

Underfitting: Too simple

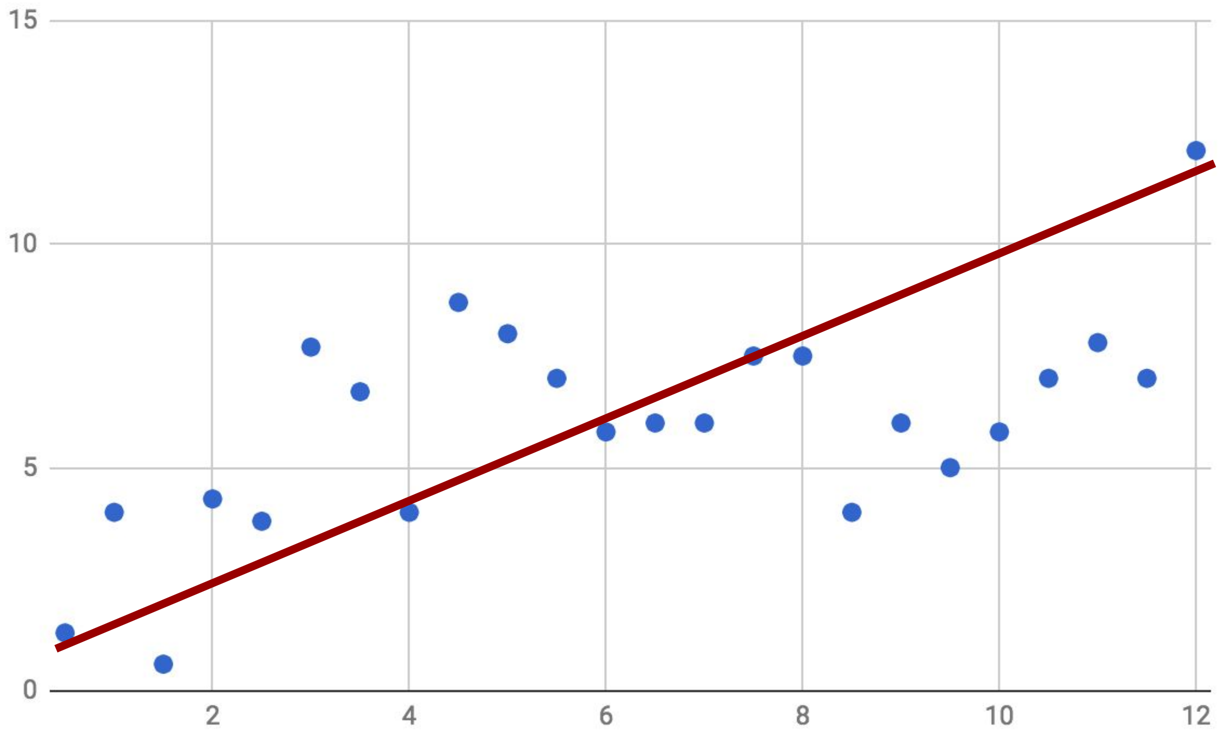

Underfitting: at least the models are consistent...

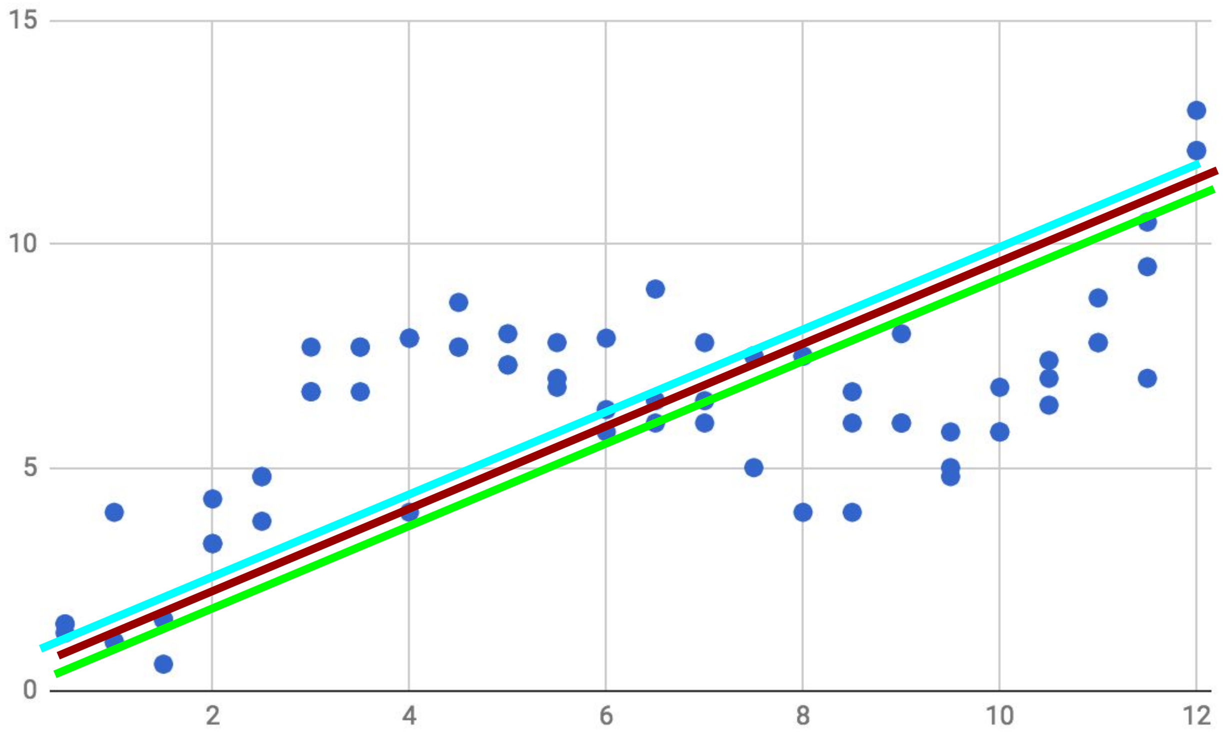

Overfitting

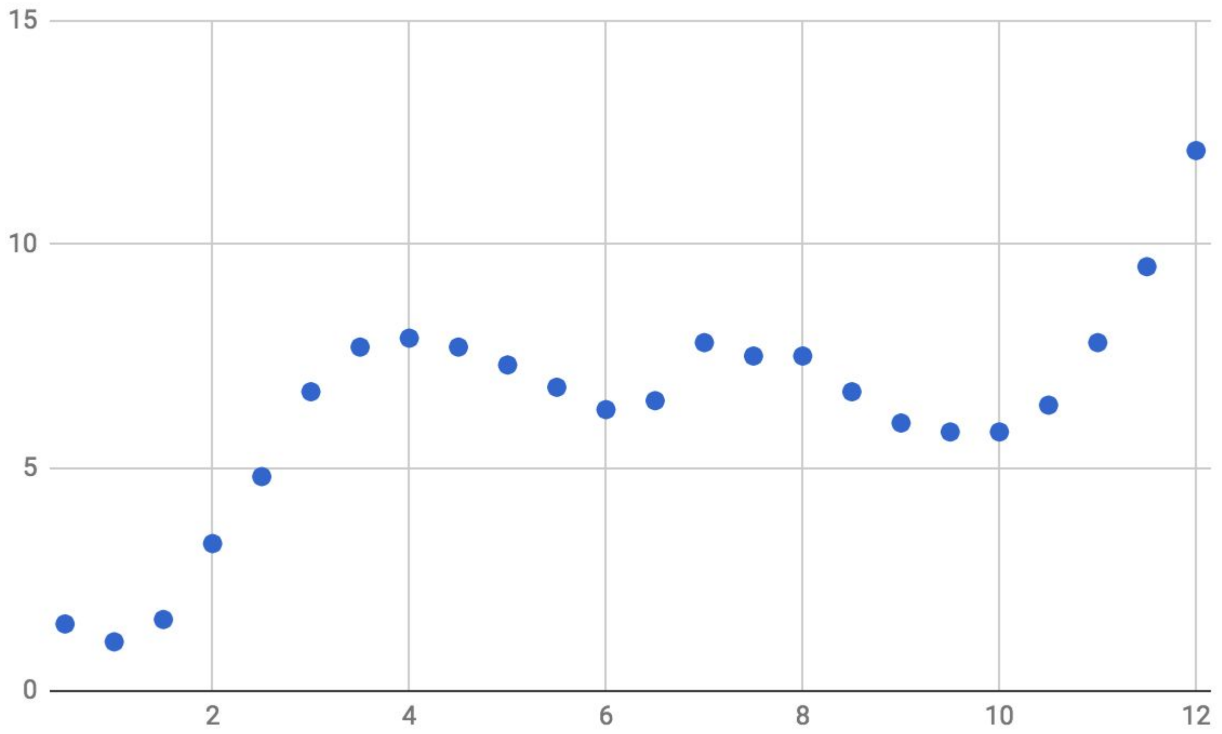

Overfitting

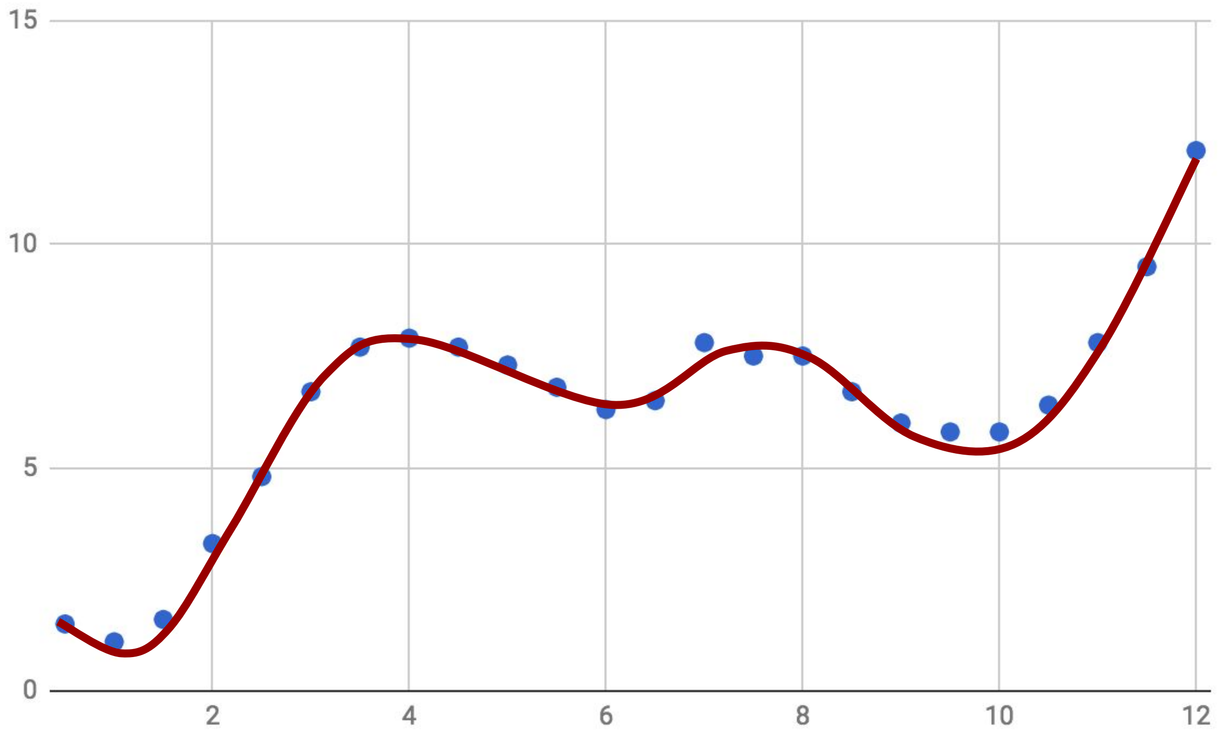

Overfitting

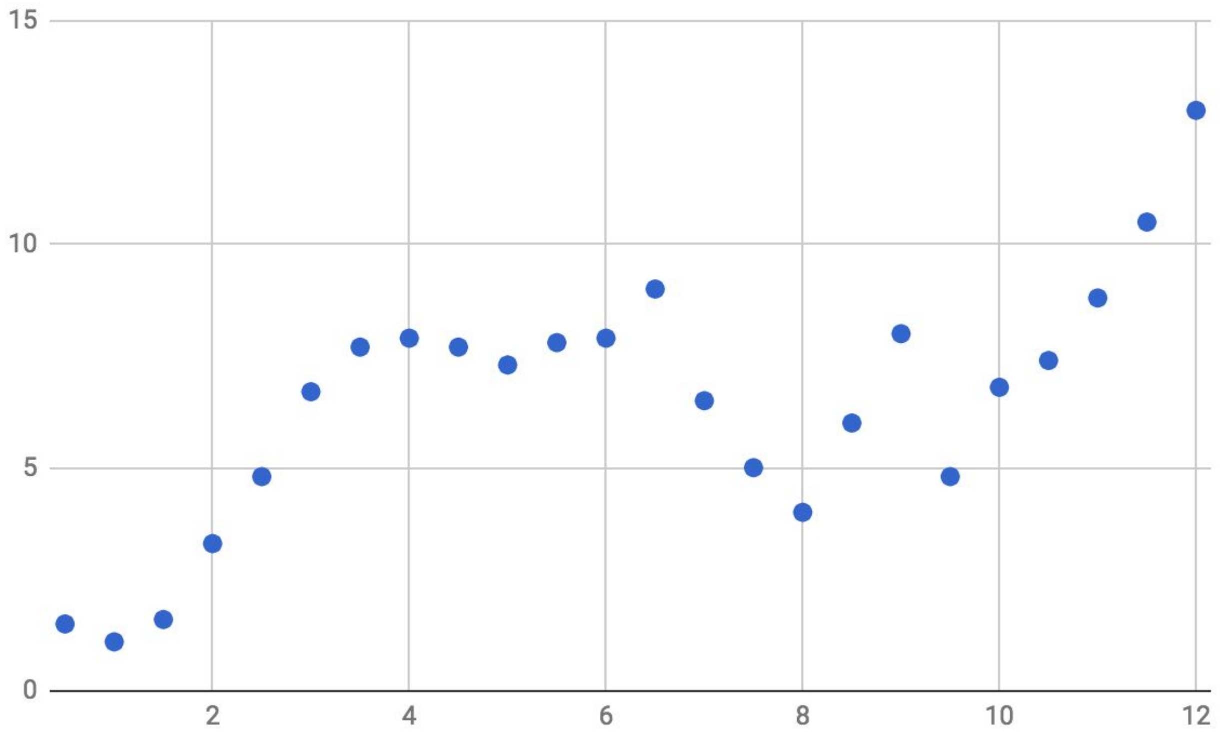

Overfitting

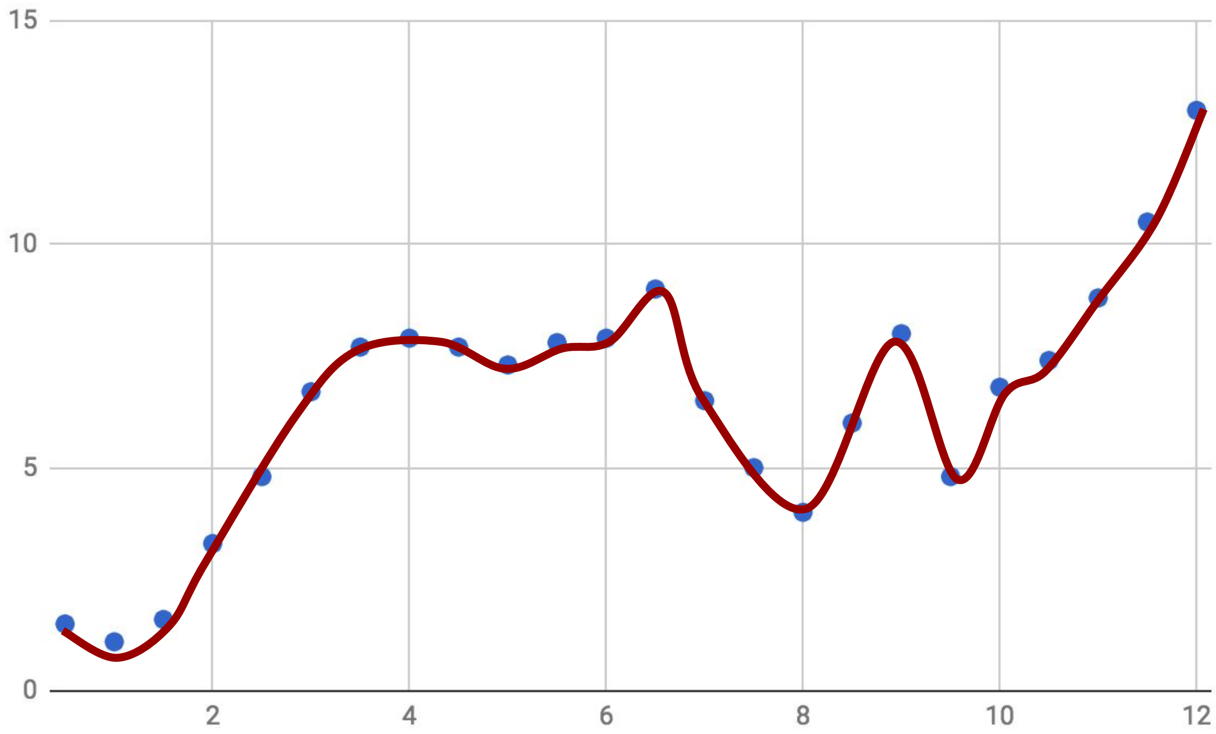

Overfitting

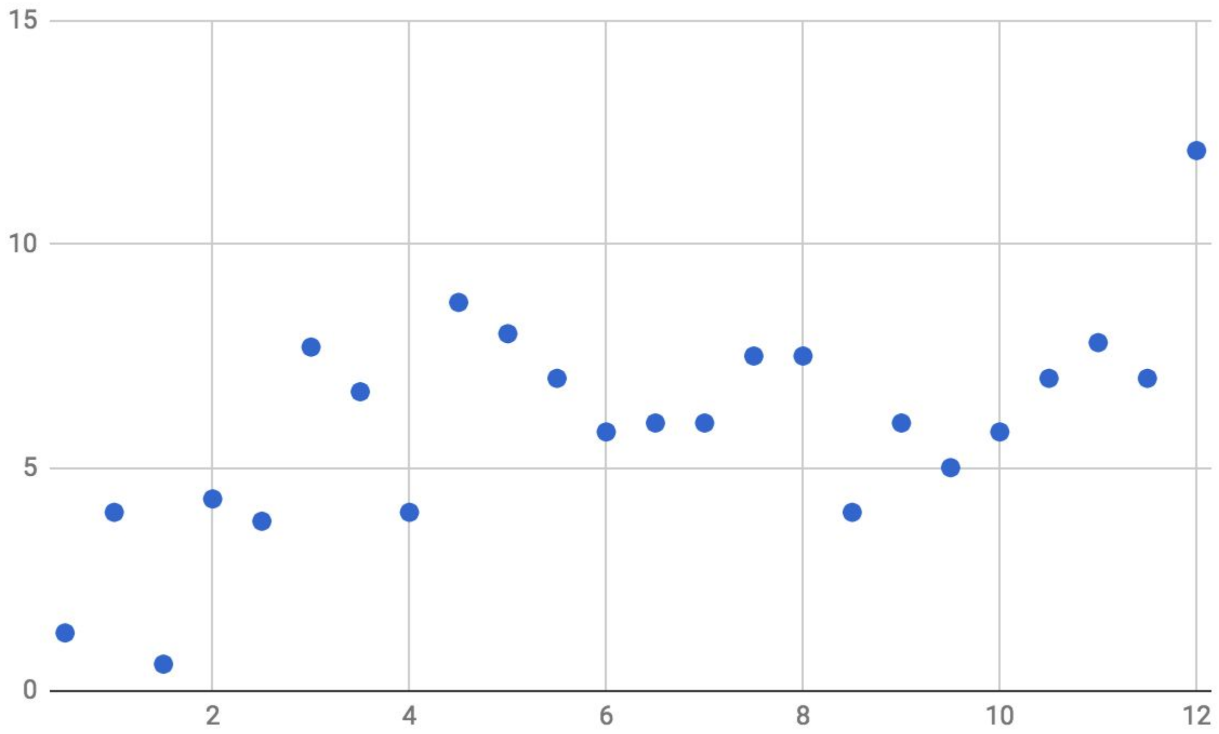

Overfitting

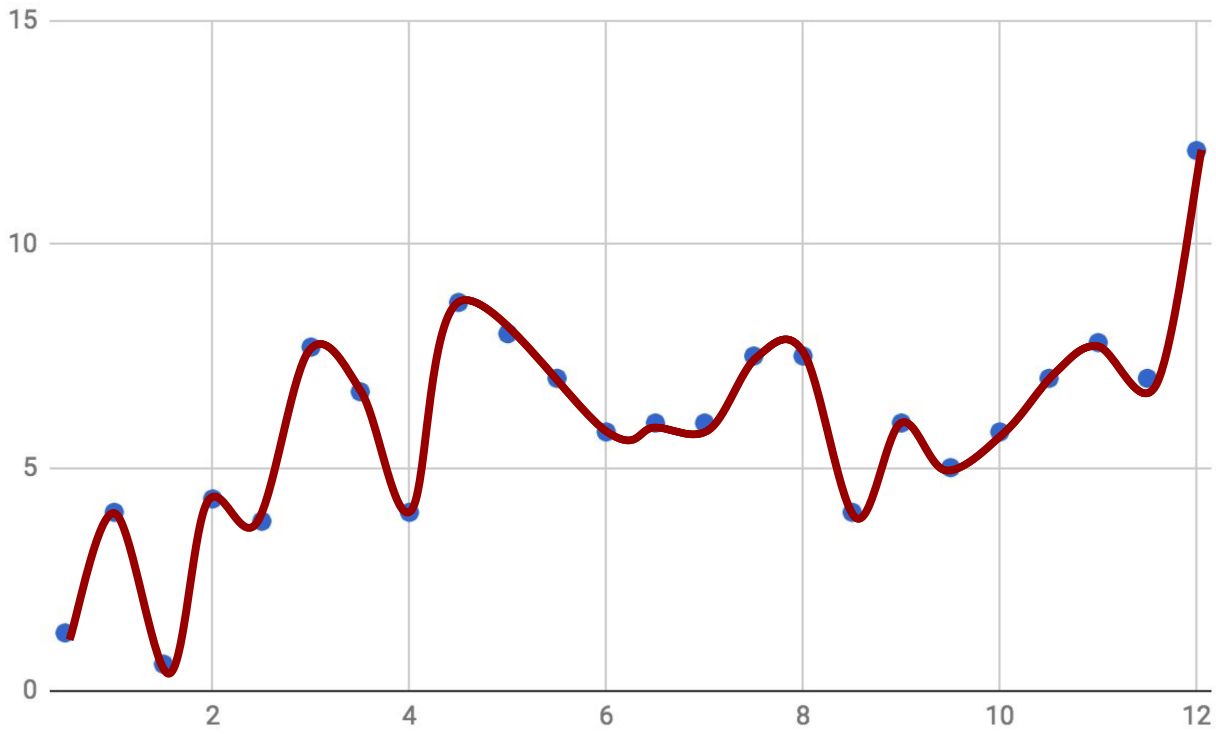

Overfitting: What's the issue?

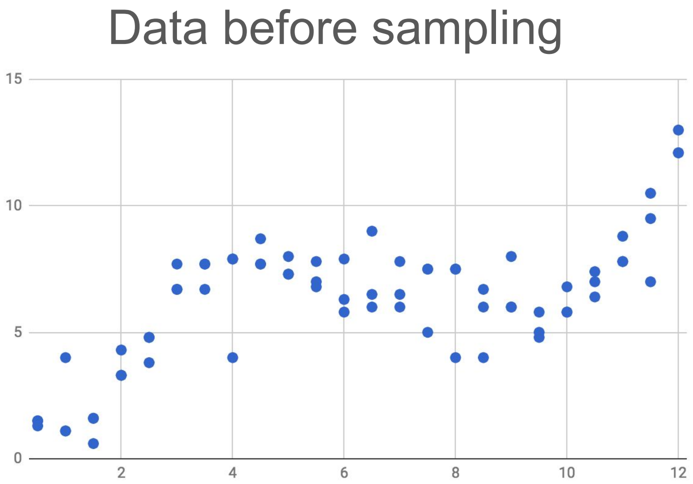

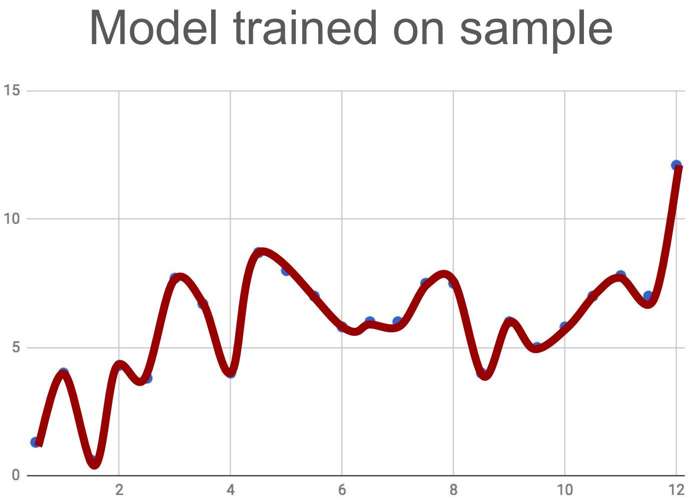

Overfitting: Inconsistent Models!

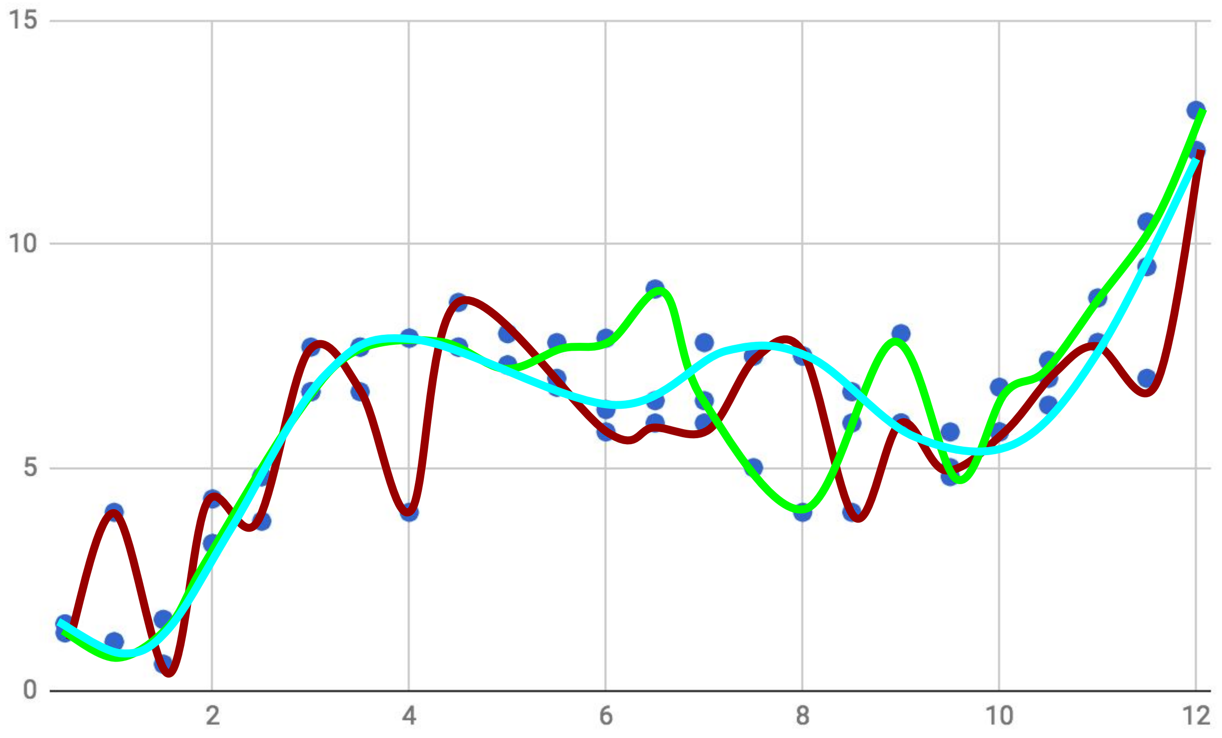

Overfitting: Results from training with high sensitivity

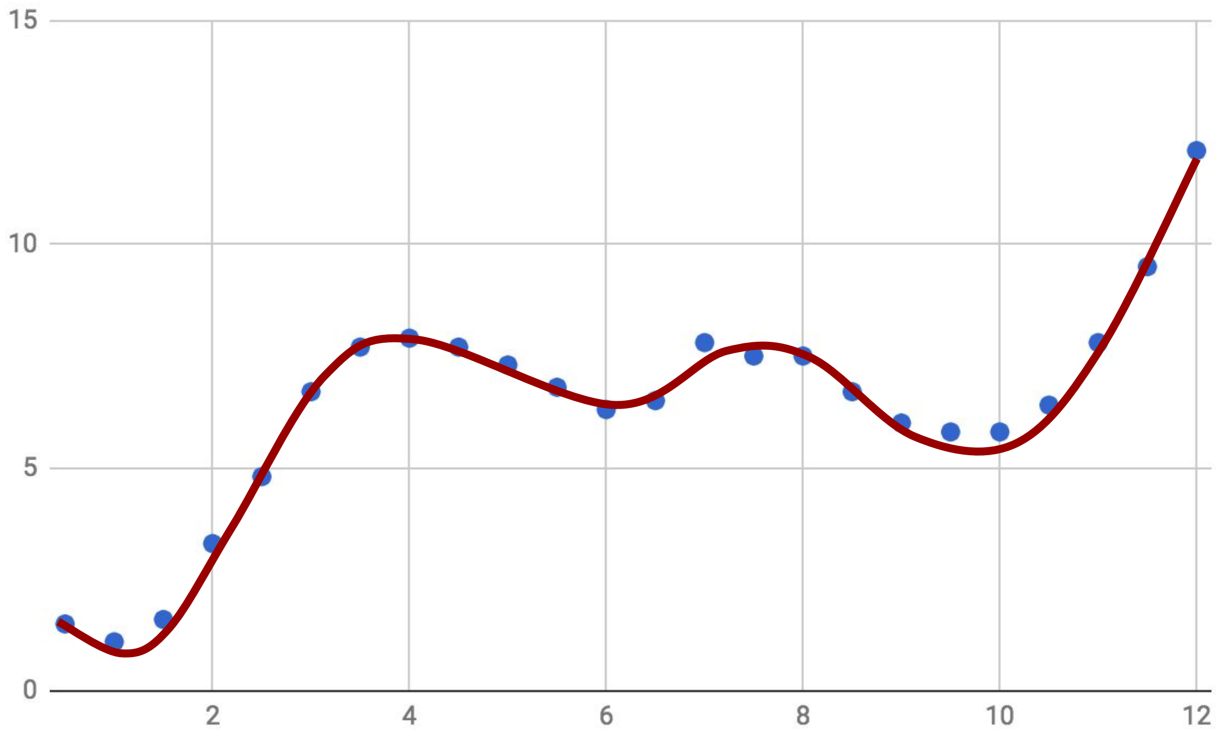

Overfitting: doesn't generalize well!

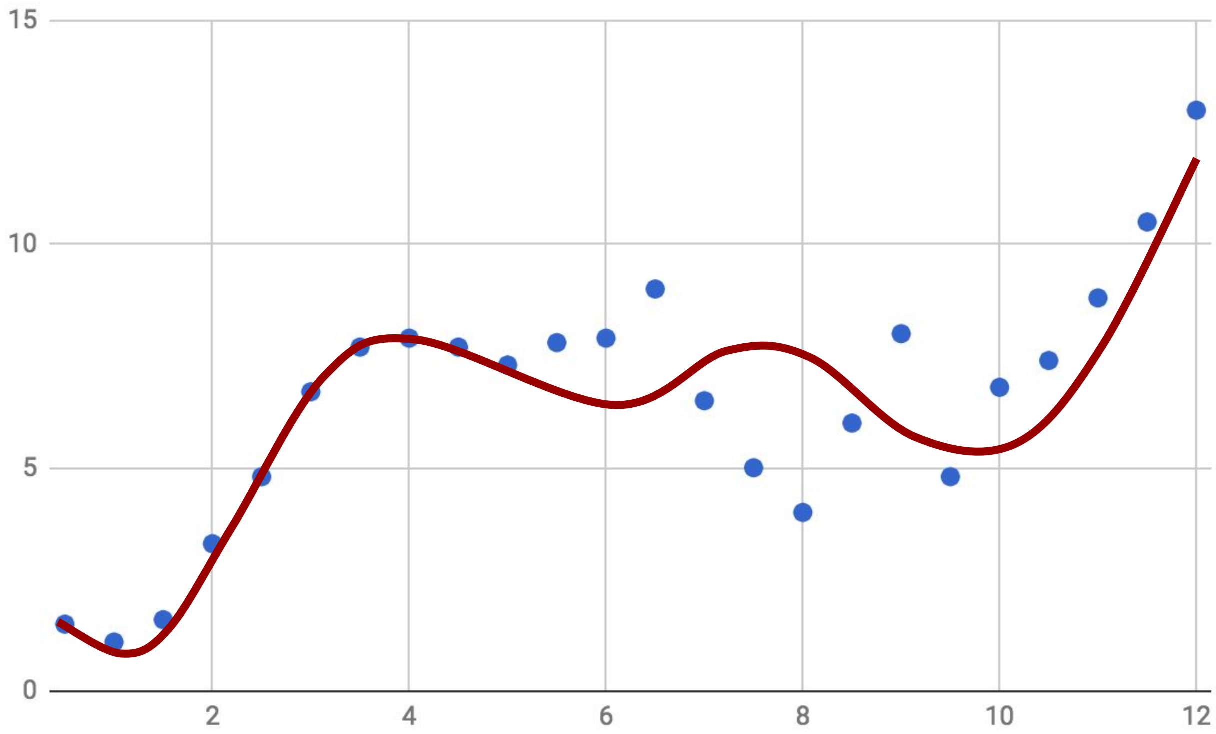

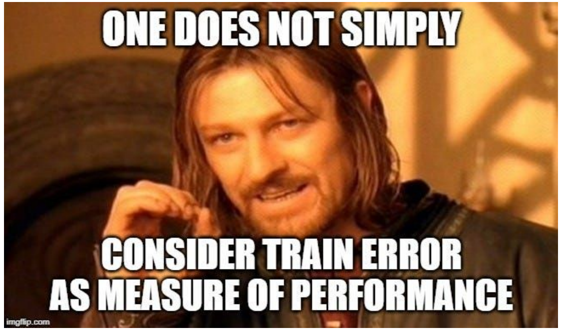

Understanding Model Error

Expected Test Error Decompositio

Framework for thinking about data:

- The world has randomness: data is randomly drawn from some distribution

- Some things have stable relations
  - Elephants are bigger than ants
  - Sun exposure can cause sun burns

→ general relation but with some variation

- Most things happen once, so we can only observe one of many the possible outcomes

Aside: how do these affect the distribution?

- Learning new things
- Practicing old things

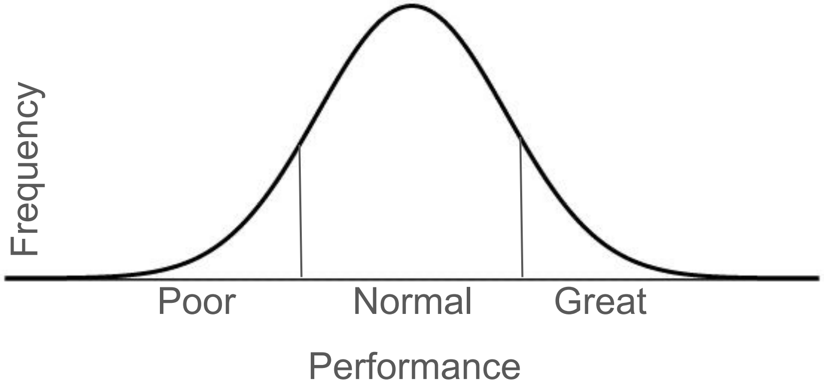

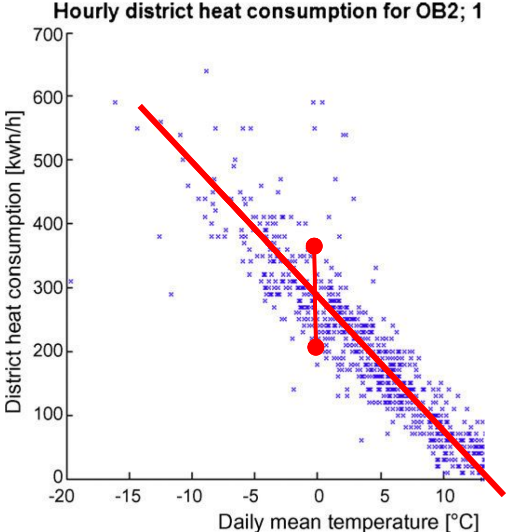

Expected Test Error Decomposition

Bias

- Error that would still exist if you had an infinite amount of training data
- Inherent to the model
  - ex. We demonstrated high bias by using a linear classifier on non-linear data

Variance

- How would your model change if you had a different training set?
- Measures how specialized your model is to your specific training set

Noise

- Measures inherent ambiguity in the data distribution
- Cannot reduce “noise” by editing algorithm

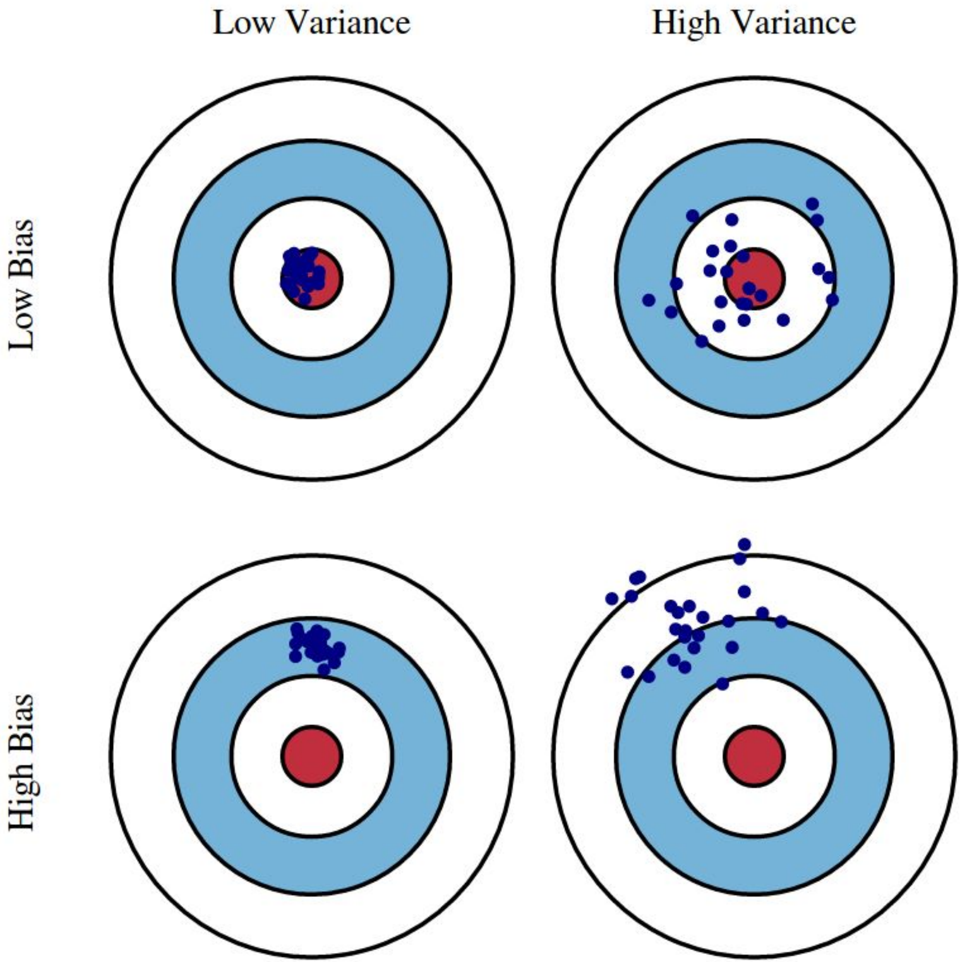

Source

What does this mean intuitively?

Bias

- Bad
- Results from incorrect assumptions in the learning algorithm

Variance

- Bad
- Results from sensitivity to fluctuations in the data

Balancing Bias and Variance

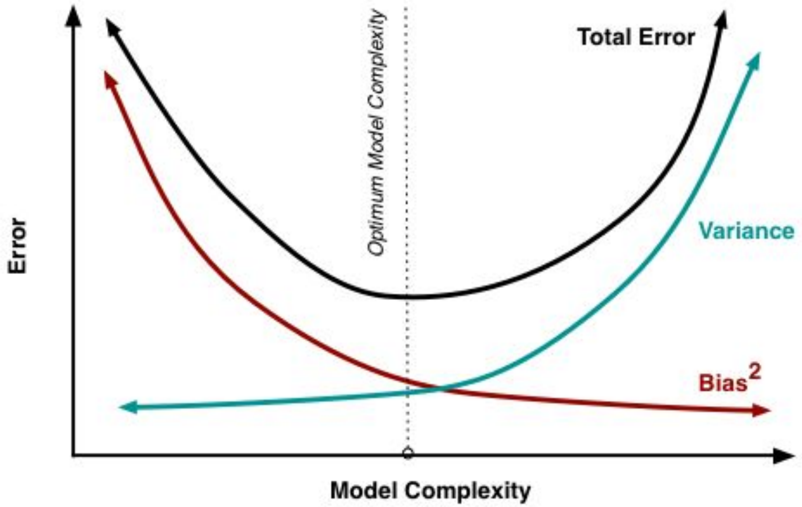

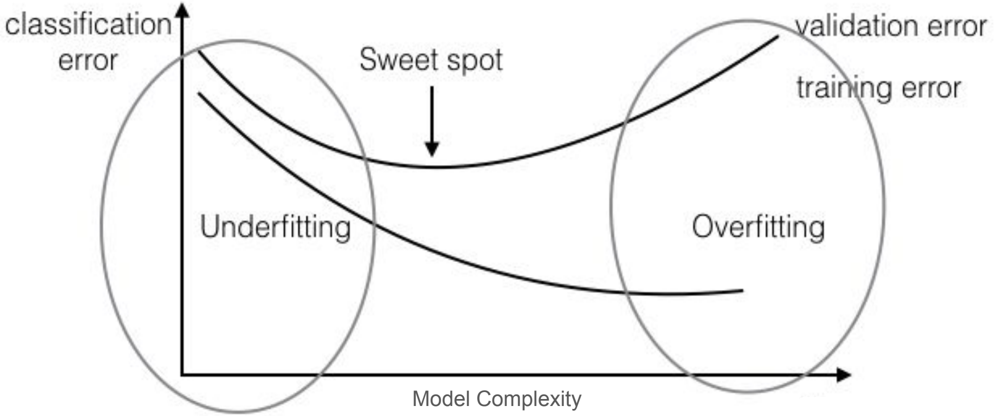

Source

Detecting and Resolving Bias and Variance

- If: High train error
  - Increase model complexity
  - Add more information (features)
  - Boost (later lecture)
  - Change model assumptions

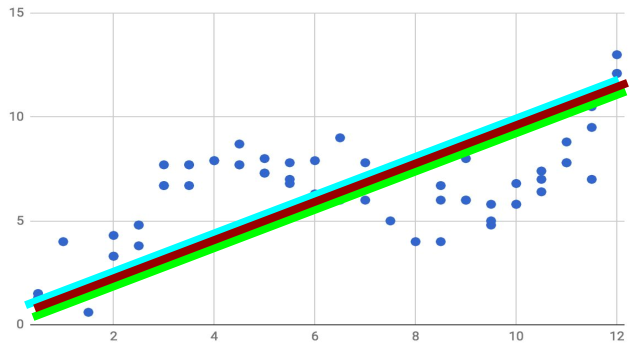

- If: Train error << test error (and test error still too high)
  - Reduce model complexity
  - Add more training data
  - Bag (later lecture)

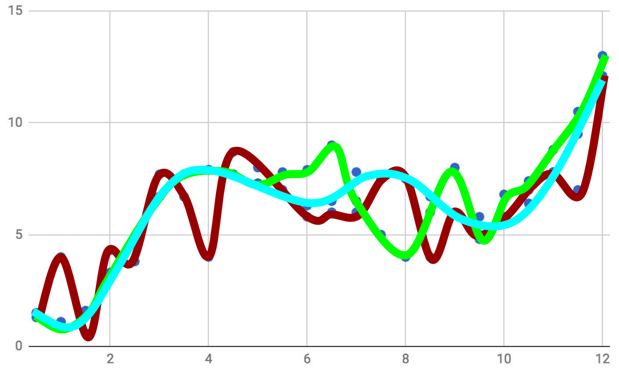

Different Topic Ahead

Any questions before we continue

Feature Selection

(adjusting models)

Methods

- **Goal:** Find subset of features that gives a good enough model, in a reasonable amount of time.
- Why:
  - More interpretable
  - More stable results
  - Less redundant/potentially misleading data
  - Faster

Correlation, r

The correlation between two variables describes to what extent changing one would change the other.

- Real-valued in [-1,1]
- A variable is always perfectly correlated with itself (correlation=1)

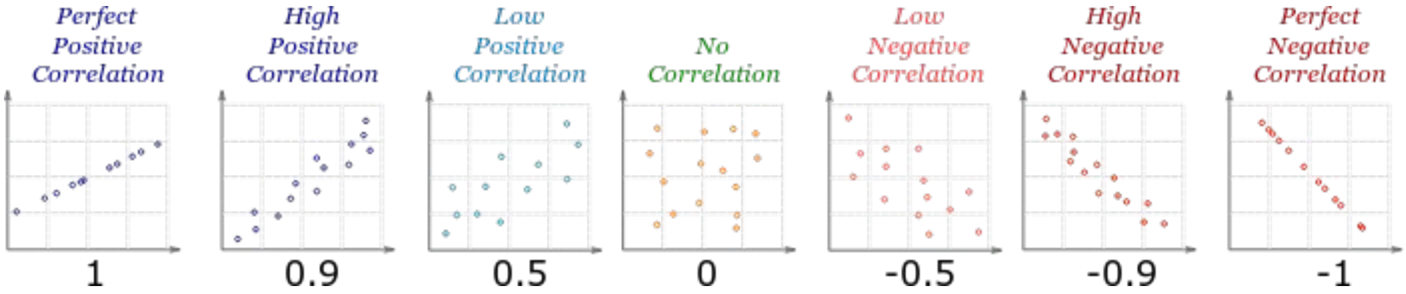

Important Case: Collinearity

Collinear: when two features have a correlation near -1 or 1

- If a feature is collinear with the target, then it's a good choice for linear regression
- If two features are collinear, they're *redundant*
    - Might as well not use one of them
    - Some models *require/assume* no collinear features
    - Takes more time, and doesn't add much information at the cost of *increased variance/sensitivity*

Demo

Final Notes

Coming Up

- **Assignment 4:** Due tonight at midnight!
- **Assignment 5:** Due midnight next Friday (10/18)
- **Mid-Semester Check-In:** Now till Wednesday (10/23)
- **Next Lecture**: Intro to Classification

Have a great Fall Break!!

CDS Education