


INFO 1998: Introduction to Machine Learning



CDS Education

We explore, learn, and educate big minds.

Lecture 6: Intro to Classifiers

INFO 1998: Introduction to Machine Learning



CDS Education

We explore, learn, and educate big minds.

Agenda

1. What is a Classifier?
2. K-Nearest Neighbors Classifier
3. Fit/Overfitting
4. Confusion Matrices



What are Classifiers?



What are Classifiers?

Classifiers are able to help answer questions like...

- “What species is this?”
- “How would consumers rate this restaurant?”
- “Which Hogwarts House do I belong to?”
- “Am I going to pass this class?”



What are Classifiers?

- Classifiers predict the class/category of a set of data points
- This class/category is based off of the target variable we are looking at
- Difference between linear regression and classifiers
 - Linear regression is used to predict the value of a **continuous variable**
 - Classifiers are used to predict **categorical or binary variables**



What are Classifiers?

Two categories of classifiers: lazy learners and eager learners

- **Lazy Learners**

- Store the training data and wait until a testing data appear
- Classification is conducted based on the most related data in the stored training data
- Less training time, more time in predicting

- **Eager Learners**

- Construct a classification model based on the given training data before receiving data for classification
- More training time, less time in predicting



K-Nearest Neighbors Classifier



What is the KNN Classifier?

- Lazy learner classifier
- Easy to interpret
- Fast to calculate
- Good for coarse analysis



How Does It Work?




KNN is based off of a simple assumption that the values **nearest** to a data point are **similar** to it

Uses the k (a user specified value) nearest data points to predict the unknown one




How Does It Work?

Most around me
got an A, maybe I
got an A as well
then.



How Does It Work? (Step-By-Step Example)


Define a k value (in this case $k = 3$)



How Does It Work? (Step-By-Step Example)

Define a k value (in this case $k = 3$)

Pick a point to predict (blue diamond)




How Does It Work? (Step-By-Step Example)

Define a k value (in this case $k = 3$)

Pick a point to predict (blue diamond)

Count the number of closest points



● 3
▲ 3




How Does It Work? (Step-By-Step Example)

Define a k value (in this case $k = 3$)

Pick a point to predict (blue diamond)

Count the number of closest points

Increase the radius until the number of points within the radius adds up to 3



● 3/3

▲ 0/3

How Does It Work? (Step-By-Step Example)


Define a k value (in this case $k = 3$)

Pick a point to predict (blue diamond)

Count the number of closest points

Increase the radius until the number of points within the radius adds up to 3

Predict the blue diamond to be a blue circle!




● 3/3

▲ 0/3



Demo



Fit/Overfitting





Overfitting

When the model corresponds too closely to training data and then isn't transferable to other data.

Can fix by:


- Splitting data into training and validation sets
- Increasing k




Relationship Between k and Fit

The k value you use has a relationship to the fit of the model

A higher k gives a smoother line, but too large of a k and it is the average of all the data



Demo



Confusion Matrix



What is a Confusion Matrix?

Table used to describe the performance of a classifier on a set of binary test data for which the true values are known

	p' (Predicted)	n' (Predicted)
p (Actual)	True Positive	False Negative
n (Actual)	False Positive	True Negative



Sensitivity

Called the **true positive rate**

Tells us how many positives are correctly identified as positives

Optimize for: airport security & initial diagnosis of fatal disease

$$\text{Sensitivity} = \text{True Positive} / (\text{True Positive} + \text{False Negative})$$



Specificity

Called the **true negative rate**

Tells us how many negatives are correctly identified as negatives

Optimize for: testing for a disease with a risky treatment

$$\text{Specificity} = \frac{\text{True Negative}}{\text{True Negative} + \text{False Positive}}$$



Question

Which is an example of when you would want **higher specificity**?

- DNA tests for a death penalty case
- Deciding which iPhone to buy
- Airport security



Overall Accuracy

Proportion of correct predictions

$$\text{Accuracy} = (\text{True Positive} + \text{True Negative}) / \text{Total}$$



Overall Error Rate

Proportion of incorrect predictions

$$\text{Error} = (\text{False Positive} + \text{False Negative}) / \text{Total}$$




Precision

Proportion of correct positive predictions among all positive predictions

$$\text{Precision} = \text{True Positive} / (\text{True Positive} + \text{False Positive})$$



Demo



Coming Up

- **Assignment 6:** Due at 4:30pm on April 14th, 2021
- **Mid-Semester Check-In:** Don't forget!! Complete before next lecture
- **Next Lecture:** Applications of Supervised Learning



CDS Education

We explore, learn, and educate big minds.